

---

# ECG Anomaly Detection and Arrhythmia Classification

---

**AUTHOR**

A.K.S.S. Wimalasena

**SUPERVISOR**

Ms. Sachini Karunaratne

Department of Mathematics

University of Ruhuna



**University of Ruhuna**

Department of Mathematics

---

# Acknowledgements

---

I sincerely thank my supervisor, Ms. Sachini Karunaratne, for her guidance, constructive comments, and encouragement throughout this research. I also thank the staff of the Department of Mathematics, University of Ruhuna, for providing the academic environment needed to complete this work.

I am grateful to the researchers and organisations that maintain PhysioNet and the MIT-BIH databases. Their carefully annotated physiological data made the experimental part of this study possible. Finally, I thank my family and friends for their patience and support during the design, implementation, testing, and writing of this thesis.

*Matara, September 2026*  
*A.K.S.S. Wimalasena*

# Abstract

---

This thesis develops a single-lead electrocardiogram (ECG) research system. A binary Extra Trees model detects abnormal beats, and a second ensemble assigns a supported abnormal subtype. Its six output codes are N, PVC, PAC, LBBB, RBBB and AFib. These are project annotation groups: AFib includes atrial flutter, and N/PVC/PAC include selected escape or supraventricular symbols rather than representing pure clinical diagnoses.

A frozen dilated U-Net autoencoder was developed from the first stored ECG channel of all 18 MIT-BIH Normal Sinus Rhythm Database subjects, with 15 fitting and three validation subjects. The anatomical identity of that channel is unspecified. The supervised models use MIT-BIH Arrhythmia Database MLII from 46 records representing 45 subject clusters. Each beat provides 223 waveform/statistical features, 15 RR features and four autoencoder residuals. The corrected pipeline filters and standardises each four-second window locally and requires observed neighbouring R peaks.

The final rerun uses the last 20% of chronologically ordered accepted beats per record for scoring, with earlier regions for model selection and fitting and a 16-beat boundary gap. The test segment had been observed during earlier development. The results therefore describe a retrospective known-patient evaluation with possible selection optimism, not confirmatory performance on new patients.

Binary accuracy was 98.81% on 20,969 beats. Conditional subtype accuracy was 97.63% on 7,958 supported abnormal beats. Whole-system accuracy was 98.07% and macro F1 was 0.9524 on 20,043 supported beats. Equal-subject accuracy was 98.03%, with a 10,000-replicate subject-clustered 95% interval of 96.29–99.25%. Its 2.95-percentage-point width meets the three-point target. However, PAC end-to-end recall was only 70.63%.

A post-fit check on complete INCART records I01 and I02, both from one source patient, used 5,416 annotated beats and a fixed threshold. Binary accuracy was 94.11%, whole-system accuracy was 92.85%, and binary precision was 64.22%. Only N and PVC had reference support. These results expose external false alarms but cannot establish population generalisation. The prototype and its reproducible evaluation are complete; class recall, independent validation, and real sensor performance remain research limitations.

# Contents

---

|                                                        |             |
|--------------------------------------------------------|-------------|
| <b>List of Figures</b>                                 | <b>viii</b> |
| <b>List of Tables</b>                                  | <b>ix</b>   |
| <b>1 Introduction</b>                                  | <b>1</b>    |
| 1.1 Introduction . . . . .                             | 1           |
| 1.2 Research Problem . . . . .                         | 4           |
| 1.3 Objective and Expected Outcomes . . . . .          | 4           |
| 1.4 Research questions . . . . .                       | 5           |
| <b>2 Literature Review</b>                             | <b>6</b>    |
| <b>3 Data</b>                                          | <b>11</b>   |
| 3.1 Data Sources and Roles . . . . .                   | 11          |
| 3.2 Data Cleaning . . . . .                            | 14          |
| 3.2.1 Cleaning Checklist . . . . .                     | 14          |
| 3.2.2 Window Construction . . . . .                    | 15          |
| 3.2.3 Cleaning Yield . . . . .                         | 16          |
| 3.2.4 Installed Use . . . . .                          | 16          |
| 3.3 Data Preparation . . . . .                         | 17          |
| 3.3.1 Prepared Objects . . . . .                       | 17          |
| 3.3.2 Splits . . . . .                                 | 18          |
| 3.3.3 Exact Model Allocation . . . . .                 | 19          |
| 3.3.4 Feeding the Frozen System . . . . .              | 20          |
| <b>4 Theoretical Foundations</b>                       | <b>21</b>   |
| 4.1 Why ECG evidence has more than one form . . . . .  | 21          |
| 4.2 Feature extraction theory . . . . .                | 22          |
| 4.3 RR timing theory . . . . .                         | 23          |
| 4.4 Normal-only reconstruction theory . . . . .        | 23          |
| 4.5 Hierarchical classification theory . . . . .       | 24          |
| 4.6 Why Extra Trees fits this feature design . . . . . | 25          |
| 4.7 Class and patient imbalance theory . . . . .       | 26          |
| 4.8 Temporal patient adaptation theory . . . . .       | 26          |
| 4.9 Validation theory . . . . .                        | 27          |
| 4.10 Theory-to-system alignment . . . . .              | 28          |

|          |                                                     |           |
|----------|-----------------------------------------------------|-----------|
| <b>5</b> | <b>Final Model Architecture</b>                     | <b>29</b> |
| 5.1      | Architecture in one sentence . . . . .              | 29        |
| 5.2      | Input signal and beat window . . . . .              | 29        |
| 5.3      | Three trained components . . . . .                  | 30        |
| 5.4      | Normal-only autoencoder . . . . .                   | 31        |
| 5.4.1    | Purpose . . . . .                                   | 31        |
| 5.4.2    | Network structure . . . . .                         | 32        |
| 5.4.3    | Normal-reference data and checkpoint . . . . .      | 33        |
| 5.4.4    | Four autoencoder residual features . . . . .        | 34        |
| 5.5      | Final 242-value feature vector . . . . .            | 35        |
| 5.6      | Binary abnormality detector . . . . .               | 36        |
| 5.7      | Abnormal-subtype classifier . . . . .               | 36        |
| 5.8      | Final decision logic . . . . .                      | 37        |
| 5.9      | Why this architecture was selected . . . . .        | 38        |
| <b>6</b> | <b>Experimental Method</b>                          | <b>39</b> |
| 6.1      | Method overview . . . . .                           | 39        |
| 6.2      | Research objective and success conditions . . . . . | 40        |
| 6.3      | Input data and feature construction . . . . .       | 40        |
| 6.4      | Temporal split . . . . .                            | 41        |
| 6.5      | Model selection and final fitting . . . . .         | 42        |
| 6.6      | Evaluation order . . . . .                          | 43        |
| 6.7      | Metrics . . . . .                                   | 44        |
| 6.8      | Subject-clustered confidence interval . . . . .     | 45        |
| 6.9      | Tests against the 90% target . . . . .              | 45        |
| 6.10     | Reproducibility and leakage controls . . . . .      | 46        |
| <b>7</b> | <b>Results and Statistical Validation</b>           | <b>47</b> |
| 7.1      | Evaluation population and supports . . . . .        | 47        |
| 7.2      | Model selection . . . . .                           | 48        |
| 7.3      | Headline temporal results . . . . .                 | 49        |
| 7.4      | Uncertainty and subject variation . . . . .         | 49        |
| 7.5      | Errors by stage and class . . . . .                 | 51        |
| 7.6      | Limited external INCART check . . . . .             | 52        |
| 7.7      | Answers to the research questions . . . . .         | 53        |
| <b>8</b> | <b>Discussion and Conclusion</b>                    | <b>55</b> |
| 8.1      | Main findings . . . . .                             | 55        |
| 8.2      | Achievement of the objectives . . . . .             | 55        |
| 8.3      | Interpretation of the corrections . . . . .         | 55        |
| 8.4      | Scientific limitations . . . . .                    | 56        |
| 8.5      | Further research . . . . .                          | 57        |
| 8.6      | Final statement . . . . .                           | 57        |

---

|                                                  |           |
|--------------------------------------------------|-----------|
| <b>Bibliography</b>                              | <b>58</b> |
| <b>A Reproducibility Record</b>                  | <b>61</b> |
| A.1 Final experiment artifacts . . . . .         | 61        |
| A.2 Frozen constants . . . . .                   | 61        |
| A.3 Saved success checks . . . . .               | 62        |
| A.4 Safety statement . . . . .                   | 62        |
| A.5 Verification and historical record . . . . . | 62        |

# List of Figures

---

|     |                                                                                                                                                                                                                                                                                                                            |    |
|-----|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----|
| 1.1 | Main components of a simplified ECG beat. Shape features describe the P–QRS–T pattern, while RR features describe timing between beats. This is an educational drawing, not a diagnostic template. . . . .                                                                                                                 | 1  |
| 1.2 | Real four-second ECG examples from the MIT–BIH Arrhythmia Database: normal record 100, PVC record 208, PAC record 232, LBBB record 109, RBBB record 118, and AFib record 203. The dashed line marks the annotated central R peak. The examples are used for explanation and are not clinical diagnostic templates. . . . . | 3  |
| 3.1 | Dataset roles and evaluation boundaries. The healthy NSRDB channel creates the frozen normal model. MITDB MLII creates the supervised feature table and is split chronologically into development and temporal-test regions. . . .                                                                                         | 12 |
| 3.2 | MLII polarity and the channel used by each model. . . . .                                                                                                                                                                                                                                                                  | 13 |
| 3.3 | Cleaning real ECG for the two roles. The healthy example is the first channel of NSRDB record 16265. The unseen-use example is MLII from MITDB record 100. Both become one 512-sample input, but their documented channel identities remain different. . . . .                                                             | 17 |
| 3.4 | Preparation of normal-reference windows and labelled MLII beats. . . . .                                                                                                                                                                                                                                                   | 20 |
| 4.1 | Two evidence scales used by the final model. Morphology describes the local shape of a beat; RR timing describes rhythm context around the beat. . . . .                                                                                                                                                                   | 21 |
| 4.2 | Feature-extraction operations behind the final 242-value vector. The tree models learn from compact numerical evidence, not from raw ECG samples directly. . . . .                                                                                                                                                         | 22 |
| 4.3 | Normal-only reconstruction principle. A larger residual can indicate novelty, but it is evidence only; the final arrhythmia decision is made by the supervised hierarchy. . . . .                                                                                                                                          | 24 |
| 4.4 | Probability logic of the final hierarchy. The subtype model is conditional: it is used only after the abnormality gate is passed. . . . .                                                                                                                                                                                  | 25 |
| 4.5 | Theoretical reason for subject-clustered validation. Clustered resampling prevents one long record from being treated as many independent people. . . . .                                                                                                                                                                  | 27 |
| 5.1 | Exact runtime order of the final ECG system. The autoencoder supplies features; the supervised tree models make the final class decision. . . . .                                                                                                                                                                          | 29 |
| 5.2 | R-peak-aligned beat window used by the final model. The R peak is at sample index 200, giving the model pre-QRS context and post-QRS recovery information. . . . .                                                                                                                                                         | 30 |
| 5.3 | Training and use of the saved three-component model. . . . .                                                                                                                                                                                                                                                               | 31 |

|     |                                                                                                                                                                                                                           |    |
|-----|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----|
| 5.4 | Autoencoder architecture with verified tensor dimensions. . . . .                                                                                                                                                         | 32 |
| 5.5 | Autoencoder fitting and validation loss. . . . .                                                                                                                                                                          | 33 |
| 5.6 | Real healthy and abnormal ECG examples through the frozen autoencoder. The residual between input and reconstruction is converted into four numerical features for the supervised detector. . . . .                       | 34 |
| 5.7 | Feature-vector layout used by both supervised models. The same 242-value vector is passed first to the abnormality detector and, if required, to the subtype classifier. . . . .                                          | 35 |
| 6.1 | Chronological model selection, refitting, and final testing. . . . .                                                                                                                                                      | 39 |
| 6.2 | Temporal split audit for every eligible MITDB MLII record. Blue beats are available for final fitting, yellow beats are the boundary gap before the test, and red beats are the held-out final 20% temporal test. . . . . | 42 |
| 6.3 | Model selection on the 64–80% validation region. The binary detector is selected by balanced accuracy and F1; the subtype classifier is selected by macro F1 and accuracy. . . . .                                        | 43 |
| 6.4 | Populations used for the three evaluation questions. . . . .                                                                                                                                                              | 44 |
| 7.1 | Three evaluation questions and their test populations. . . . .                                                                                                                                                            | 47 |
| 7.2 | Reference-class composition of the corrected temporal test. . . . .                                                                                                                                                       | 48 |
| 7.3 | Pooled temporal scores after the preprocessing correction. . . . .                                                                                                                                                        | 49 |
| 7.4 | Subject-clustered confidence intervals. . . . .                                                                                                                                                                           | 50 |
| 7.5 | Corrected whole-system accuracy across 45 known-patient clusters. . . . .                                                                                                                                                 | 50 |
| 7.6 | True and correctly classified supported abnormal beats. PAC remains below the 90% recall target. . . . .                                                                                                                  | 51 |
| 7.7 | Whole-system confusion matrix for 20,043 supported temporal beats. . . . .                                                                                                                                                | 52 |

# List of Tables

---

|     |                                                                                                                     |    |
|-----|---------------------------------------------------------------------------------------------------------------------|----|
| 1.1 | Normal rhythm and the five arrhythmia types considered in this study. . . . .                                       | 2  |
| 2.1 | Main papers referenced in this study and their contribution. . . . .                                                | 9  |
| 3.1 | Data sources and their strictly separated roles. . . . .                                                            | 11 |
| 3.2 | Mapping from MITDB expert annotations to project outputs. . . . .                                                   | 14 |
| 3.3 | Cleaning checklist applied before a signal can become a model input. . . . .                                        | 15 |
| 3.4 | Training, validation, and temporal-test data supplied to each model. . . . .                                        | 19 |
| 4.1 | How the theory maps to the implemented final system. . . . .                                                        | 28 |
| 5.1 | Organised feature groups in the final architecture. . . . .                                                         | 35 |
| 5.2 | Reason for each major final-architecture decision. . . . .                                                          | 38 |
| 6.1 | Exact evidence supplied to the final supervised models. . . . .                                                     | 40 |
| 7.1 | Distinct denominators for the three evaluation questions. . . . .                                                   | 48 |
| 7.2 | Corrected temporal results. Confidence intervals estimate equal-subject accuracy, not pooled beat accuracy. . . . . | 49 |
| 7.3 | Whole-system performance by supported abnormal output group. . . . .                                                | 51 |
| 7.4 | Corrected model on complete INCART I01 and I02 recordings. . . . .                                                  | 53 |
| A.1 | Settings used for the final experiment. . . . .                                                                     | 61 |


# 1 Introduction

---

## 1.1 Introduction

An electrocardiogram (ECG) records the small voltage changes that occur when electrical activity moves through the heart. Electrodes on the body see this activity from different directions. In an ECG trace, the horizontal axis shows time and the vertical axis shows voltage. ECG analysis is useful because the pattern changes when a heartbeat starts in a different place, happens at a different time, or travels through the heart in a different way.

**Components of one ECG beat.** A normal heartbeat starts when the sinoatrial node creates an impulse. The P wave shows atrial depolarisation. The short, sharp QRS complex mainly shows ventricular depolarisation. The T wave shows ventricular repolarisation. The time between two neighbouring R peaks is called the RR interval. It gives information about heart rate and rhythm regularity [1]. Figure 1.1 shows these parts in a simple way. In real recordings, the same beat can look inverted, stretched, or noisy because of the lead, the patient, the electrodes, and the recording conditions.



**Figure 1.1:** Main components of a simplified ECG beat. Shape features describe the P–QRS–T pattern, while RR features describe timing between beats. This is an educational drawing, not a diagnostic template.

**Arrhythmia and the classes considered.** An arrhythmia is a problem in where the heart's electrical activity starts, when it happens, how it is ordered, or how it travels. It may affect one beat, many repeated beats, or a longer rhythm. This study uses one normal output and five arrhythmia outputs. Table 1.1 summarises the main physiological cause and visible ECG behaviour for each one. These descriptions define the model task only. A clinical diagnosis still needs the full ECG and patient context [2, 3].


**Table 1.1:** Normal rhythm and the five arrhythmia types considered in this study.

| Type                    | Main reason                                                                                                 | Typical ECG behaviour and importance                                                                                                                                                           |
|-------------------------|-------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Normal sinus rhythm (N) | The impulse begins in the sinoatrial node and follows the normal conduction path.                           | Repeated P–QRS–T patterns and generally regular RR intervals. It supplies the reference against which unusual beats are compared.                                                              |
| PVC                     | A premature impulse begins in a ventricle instead of following the normal path from the atria.              | The beat occurs early and often has a wide, unusual QRS, followed by a pause. Detection is important because it is a clear ventricular abnormality and can occur repeatedly.                   |
| PAC                     | A premature impulse begins in an atrial location outside the sinoatrial node.                               | The beat occurs early; the P wave may change, while the QRS can remain narrow and similar to normal. Its subtle shape makes PAC difficult to separate from normal beats.                       |
| LBBB                    | Electrical conduction through the left bundle branch is delayed or blocked.                                 | The ventricles activate in an altered sequence, producing a broad, changed QRS. Recognition is important because the pattern can be associated with underlying cardiac disease.                |
| RBBB                    | Electrical conduction through the right bundle branch is delayed or blocked.                                | The later part of the QRS changes and can show a second peak in some leads. It must be separated from LBBB because the delayed side is different.                                              |
| AFib                    | Atrial electrical activity becomes rapid and disorganised rather than following one regular atrial impulse. | Consistent P waves disappear and RR intervals become irregular. Identification is important because AFib is linked with clinically important complications and requires rhythm-level evidence. |

Figure 1.2 shows real labelled ECG segments from the MIT–BIH Arrhythmia Database instead of drawings. Each panel comes from the fixed final 20% of its record. For each stated record and class, the displayed window is the one closest to that held-out group’s mean waveform by squared distance. This rule gives a representative example without choosing the most dramatic-looking beat. The signals were resampled to 128 Hz, standardised, and aligned

so the labelled R peak is at 1.5625 seconds.

### Real labeled ECG segments from the locked final 20%



**Figure 1.2:** Real four-second ECG examples from the MIT-BIH Arrhythmia Database: normal record 100, PVC record 208, PAC record 232, LBBB record 109, RBBB record 118, and AFib record 203. The dashed line marks the annotated central R peak. The examples are used for explanation and are not clinical diagnostic templates.

These examples show why one simple threshold is not enough. PVC, LBBB, and RBBB mainly change beat shape. PAC can occur early while still looking close to a normal beat. AFib is often clearer from changing intervals and rhythm context. For this reason, the final system combines MLII waveform shape, frequency information, RR history, and disagreement with a learned normal reconstruction.

Finding these patterns matters because long ECG recordings can contain too many beats for continuous manual review. An automated research system can mark suspicious regions, count events, and help a reviewer focus attention. At the same time, an error can miss an abnormality or create an unnecessary alarm. Therefore, accuracy, per-class recall, confidence intervals, and clear limits are all needed. The system developed here is for research evaluation only. It is not a replacement for a clinician.

## 1.2 Research Problem

The central research problem is stated below.

**Main research problem:** How can an explainable two-stage ECG system use normal-only reconstruction, beat morphology, and heartbeat timing to detect abnormal beats and correctly identify PVC, PAC, LBBB, RBBB, and AFib with more than 90% whole-system accuracy on a fixed temporal test, while providing statistically valid evidence across patients?

The main problem contains four connected sub-problems:

1. ECG shape varies between patients and leads, so harmless personal morphology can resemble an abnormality.
2. A two-stage system can fail at the abnormality gate or at subtype classification; a missed abnormal beat never reaches the second stage.
3. Normal and arrhythmia classes have unequal numbers of beats and patients, so ordinary accuracy can hide failure on a rare class.
4. Thousands of neighbouring beats from one patient are correlated. Testing must keep the final signal unseen and must calculate uncertainty using complete patient clusters rather than treating beats as independent people.

The final experiment handles these problems with a chronological split. The first 64% of each MITDB record trains candidate models, the next 16% selects the models and threshold, and the final 20% is reserved for temporal testing. The present model fitting excludes test rows, but the segment was observed during earlier development. Percentages count accepted beats in chronological order, rather than exact recording duration. The result is therefore a retrospective patient-adapted temporal evaluation, with possible selection optimism and no confirmatory claim about completely new patients.

## 1.3 Objective and Expected Outcomes

**Main objective.** The main objective is to develop and statistically evaluate an ECG system that first detects abnormal beats and then identifies five supported arrhythmia types in later unseen signal from known patients.

**Specific objectives.**

1. Train a normal-only autoencoder with the complete first ECG channel from all 18 MIT-BIH Normal Sinus Rhythm Database subjects and use its reconstruction errors as evidence.
2. Use the earlier 80% of every MITDB record containing MLII for supervised development while keeping every final 20% portion unavailable.

3. Build a binary abnormality detector using ECG morphology, spectrum, RR timing, and autoencoder-residual features;
4. Build a second classifier for PVC, PAC, LBBB, RBBB, and AFib using the same MLII signal and RR context with a next-beat delay;
5. measure abnormality-detection accuracy, conditional subtype accuracy, and complete six-class system accuracy separately;
6. report the true and correctly classified count for every arrhythmia; and
7. validate uncertainty with equal-subject accuracy, a complete-subject bootstrap confidence interval, and subject-level statistical tests.

**Expected outcomes.** The planned outcomes are:

- a reproducible, deployable two-stage research prototype;
- abnormality-detection, subtype, and whole-system point accuracies above 90% on the fixed temporal test;
- a whole-system 95% confidence interval whose lower bound exceeds 90% and whose total width is no more than three percentage points;
- per-arrhythmia results that reveal whether every supported class reaches the accuracy target; and
- an honest statement of the data, statistical, hardware, and clinical limits of the model.

These are evaluation targets, not assumed conclusions. Chapter 7 reports which targets were reached and which were not.

## 1.4 Research questions

- RQ1** How does the frozen NSRDB reconstruction branch contribute evidence to the supervised hierarchy, and has its independent benefit been measured?
- RQ2** What are binary, conditional-subtype, and complete six-class performance on the retrospective temporal evaluation?
- RQ3** What do equal-subject accuracy and subject-clustered uncertainty show about variation across known patients?
- RQ4** Do the measured accuracy, interval precision, and class recalls meet the study targets, and what evidence exists for external use?

The labels PVC, PAC, N and AFib are operational annotation groups defined in Chapter 3; AFib includes flutter. A successful software implementation is distinct from meeting every statistical target.

## 2 Literature Review

---

This chapter reviews the work most relevant to the thesis questions: can a system detect and name ECG abnormalities, and how does performance change when earlier ECG from the same patient is available? The review focuses on four areas: ECG datasets and evaluation protocols, supervised heartbeat classification, normal-only anomaly detection, and statistical validation. It is a focused narrative review, not a systematic review or meta-analysis. Papers were selected because they explain a design choice, show an important alternative, or reveal a risk in patient-independent evaluation.

Headline accuracy is not used to rank the papers. ECG studies often use different leads, labels, record lengths, patient splits, and output units. For example, a patient-specific beat classifier, a 12-lead diagnostic model, and the six-class system in this thesis answer different questions. Their accuracy values cannot be compared directly. This matters because a high score may come from an easier class mapping, or from placing beats from the same person in both training and test data.

PhysioNet provides open physiological databases and common software tools [4]. The MIT-BIH Arrhythmia Database (MITDB) became a widely used benchmark because it contains expert beat annotations and two ECG channels [2, 5]. The MIT-BIH Normal Sinus Rhythm Database provides 18 real long-term recordings from subjects without significant arrhythmias [6]. It is suitable for normal-only learning, but its headers identify the channels only as ECG1 and ECG2. The present study therefore uses one fixed first channel and reports this uncertainty instead of assigning an unsupported anatomical lead name.

The evaluation protocol is as important as the database. A beat-wise random split can place very similar beats from one patient in both training and test sets. An inter-patient split instead tests whether the learned boundary transfers to a different person. De Chazal et al. established an influential inter-patient heartbeat study using ECG morphology and heartbeat-interval features [3]. Later research continued to treat patient variation as a central problem rather than assuming that all beats are independent [7, 8].

This evidence supports a clearly defined evaluation protocol instead of a single ambiguous score. The final experiment keeps chronological order, trains on the first 80% of every record, and tests on the final 20%. The boundaries count accepted beats. Although the rerun preserves chronology, earlier development already exposed the final segment. Its results are retrospective temporal evidence with possible selection optimism. The separate NSRDB cohort is used only for the normal autoencoder; no healthy record trains either supervised tree model.

Early automatic ECG systems used measured morphology and timing. De Chazal et al. combined waveform information with RR intervals, showing why beat shape and beat timing should be treated as complementary evidence [3]. This idea remains useful even when the classifier changes: morphology helps describe PVC and bundle-branch beats, while RR history

helps describe premature beats and irregular rhythms.

Huang et al. studied inter-patient separation of normal, LBBB, and RBBB beats with ensemble classifiers [9]. That study is particularly relevant because bundle-branch morphology can persist through a record. However, its three-class task is narrower than the present N/PVC/PAC/LBBB/RBBB/AFib task. Its results therefore justify using morphology and ensemble learning, but they cannot be copied as an expected accuracy for this thesis.

Feature-based models have two practical advantages for this dataset. They make the evidence easier to inspect, and they can be controlled when only a few patients contain a class. Their weakness is that hand-designed summaries may miss useful waveform detail. This trade-off led to the combined design used here: pooled waveform and RR features are kept, while an autoencoder adds learned normal-reconstruction evidence.

One-dimensional convolutional neural networks can learn ECG morphology directly. Kiranyaz et al. developed a patient-specific real-time classifier [10]. Its patient adaptation is useful when labelled beats from the target person exist. It is therefore directly relevant to the final patient-adapted setting, but it does not answer performance before seeing a new patient. Acharya et al. classified arrhythmias from ECG segments with a convolutional network [11]. Segment-level work demonstrates the value of learned temporal features, but its label definition and split must match the target study before accuracy can be compared.

Large deep-learning studies have also reported strong performance on ambulatory single-lead recordings and 12-lead diagnostic ECGs [12, 13]. These works show the potential of deep networks when many patients and labels are available. They also show why dataset scale matters. Their databases, lead systems, and diagnostic outputs are different from the small six-class MITDB experiment, so they are evidence for a method family rather than a performance baseline for this thesis.

Sequence models use neighbouring beats to add rhythm context. Mousavi et al. compared intra-patient and inter-patient heartbeat classification with a sequence-to-sequence approach [7]. Zhang et al. added an adversarial subject classifier so that the learned representation would contain less patient identity [8]. These approaches address generalisation directly, but they also need enough independent patients. The present study has only 45 eligible subject clusters, so the final system uses a simpler tree hierarchy. This is a data-availability decision, not a claim that deep learning is generally inferior.

Autoencoders learn a compressed representation by reconstructing their input [14]. Denoising training corrupts the input but retains a clean target, encouraging the network to learn stable structure rather than an exact copy [15]. In anomaly detection, the model is often trained on normal sequences and reconstruction error is used as evidence that a new sequence differs from the learned normal pattern [16, 17].

Temporal convolutional autoencoders use convolution and dilation to cover longer time patterns efficiently. Thill et al. evaluated this idea on time-series anomaly tasks that included ECG [18]. ECG-specific studies have also investigated convolutional variational and adversarial autoencoders [19, 20]. Together, these studies support using a convolutional autoencoder for normal ECG representation.

The literature does not mean that reconstruction error is a diagnosis. A high-capacity model may reconstruct an abnormal input well, while noise or a different electrode position may

increase the error of a normal input. For that reason, this thesis does not use one reconstruction threshold as the final answer. It uses four residual summaries as features in a supervised abnormal detector. The normal-only idea is kept, but the final decision can combine residual, morphology, spectral, and RR evidence.

ECG datasets contain unequal class counts and unequal numbers of patients per class. A high overall accuracy can therefore hide failure on a rare arrhythmia. Precision–recall analysis is more informative than ROC analysis in strongly imbalanced binary problems [21], while class-balanced learning can increase the influence of rare classes [22]. This thesis therefore reports sensitivity, specificity, balanced accuracy, macro F1, and per-class recall in addition to accuracy.

Thousands of beats from a small number of subjects do not become thousands of independent patients. If individual beats are resampled, the confidence intervals become too narrow. Patient-cluster bootstrap methods keep within-patient dependence [23]. The final experiment gives each of 45 subject clusters equal weight and resamples complete subjects. This also follows modern guidance to report the data split, model development, and validation population clearly [24].

The literature provides strong individual ideas: normal-only reconstruction, learned waveform features, morphology and RR features, patient-invariant learning, specialised bundle-branch classification, and clustered statistical validation. However, reported results usually address only part of the present problem. They may use a different class mapping, adapt to the test patient, use more leads, or evaluate anomaly detection without naming the abnormal type.

The gap addressed by this thesis is therefore not simply the need for another classifier. The need is to evaluate a complete two-stage system while making the generalisation claim explicit:

1. learn normal reconstruction from one fixed channel of all 18 real MIT–BIH normal-sinus recordings;
2. use a chronological 80/20 experiment for patient-adapted temporal evidence;
3. keep each final test portion unavailable during selection; and
4. report separate abnormal, subtype, whole-system, patient-level, and uncertainty results.

Table 2.1 summarises the papers that contributed most directly to the data protocol, model architecture, and evaluation method used in this research.

**Table 2.1:** Main papers referenced in this study and their contribution.

| Paper title                                                                                                                       | Author(s)                                                                   | Main contribution used in this research                                                                                                                                                                  |
|-----------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| MIT-BIH Normal Sinus Rhythm Database [6]                                                                                          | George B. Moody                                                             | Provides the 18 real long-term recordings without significant arrhythmias used for normal-only autoencoder development; its generic channel names also motivate explicit reporting of lead uncertainty.  |
| Automatic Classification of Heartbeats Using ECG Morphology and Heartbeat Interval Features [3]                                   | Philip de Chazal, Maria O'Dwyer, and Richard B. Reilly                      | Established the value of combining ECG morphology with RR intervals and carefully defining the evaluation population. This supports the morphology, timing, and reporting choices of the present system. |
| Reducing the Dimensionality of Data with Neural Networks [14]                                                                     | Geoffrey E. Hinton and Ruslan R. Salakhutdinov                              | Provides the representation-learning basis of the autoencoder used to learn normal ECG structure.                                                                                                        |
| Extracting and Composing Robust Features with Denoising Autoencoders [15]                                                         | Pascal Vincent, Hugo Larochelle, Yoshua Bengio, and Pierre-Antoine Manzagol | Supports adding noise to the input while keeping a clean reconstruction target, so the autoencoder learns stable normal patterns.                                                                        |
| Temporal Convolutional Autoencoder for Unsupervised Anomaly Detection in Time Series [18]                                         | Markus Thill, Wolfgang Konen, Hao Wang, and Thomas Baeck                    | Supports temporal convolution, dilation, and reconstruction error for detecting unusual time-series patterns, including ECG signals.                                                                     |
| Inter- and Intra-patient ECG Heartbeat Classification for Arrhythmia Detection: A Sequence to Sequence Deep Learning Approach [7] | Sajad Mousavi and Fatemeh Afghah                                            | Shows the importance of neighbouring-beat rhythm context and clearly separates intra-patient from inter-patient evaluation.                                                                              |
| Interpatient ECG Heartbeat Classification with an Adversarial Convolutional Neural Network [8]                                    | Jing Zhang, Aiping Liu, Deng Liang, Xun Chen, and Min Gao                   | Explains patient-identity leakage and why results from known-patient temporal tests must not be presented as new-patient performance.                                                                    |
| Detection of Inter-Patient Left and Right Bundle Branch Block Heartbeats in ECG Using Ensemble Classifiers [9]                    | Huifang Huang, Jie Liu, Qiang Zhu, Rui Wang, and Guangshu Hu                | Supports morphology and ensemble classifiers for LBBB and RBBB, while showing that this specialised three-class task is narrower than the thesis task.                                                   |
| The Precision–Recall Plot Is More Informative than the ROC Plot When Evaluating Binary Classifiers on Imbalanced Datasets [21]    | Takaya Saito and Marc Rehmsmeier                                            | Justifies reporting precision, recall, and F1 because the ECG classes are imbalanced and accuracy alone can hide failure.                                                                                |

| Paper title                                                                                                                         | Author(s)              | Main contribution used in this research                                                                                             |
|-------------------------------------------------------------------------------------------------------------------------------------|------------------------|-------------------------------------------------------------------------------------------------------------------------------------|
| Bootstrap Estimation of Diagnostic Accuracy with Patient-Clustered Data [23]                                                        | Carolyn M. Rutter      | Provides the reason for resampling complete patient records instead of treating correlated beats as independent observations.       |
| TRIPOD+AI Statement: Updated Guidance for Reporting Clinical Prediction Models That Use Regression or Machine Learning Methods [24] | Gary S. Collins et al. | Supports transparent reporting of the patient split, model-development process, temporal evaluation, limitations, and intended use. |

# 3 Data

This chapter identifies the ECG sources, explains how the real signals are cleaned, and shows how examples enter model fitting, validation, and the temporal test. Keeping these roles separate prevents healthy data or test labels from being used in the wrong part of the study.

## 3.1 Data Sources and Roles


Two public PhysioNet databases supply model development data. A separate, limited external INCART check is reported in the results chapter. The MIT–BIH Normal Sinus Rhythm Database (NSRDB) supplies only healthy ECG for the autoencoder. The MIT–BIH Arrhythmia Database (MITDB) supplies labeled normal and abnormal beats for the detector, subtype classifier, and final temporal evaluation [4, 6, 5].

**Table 3.1:** Data sources and their strictly separated roles.

| Source                                          | Available data                                                                                                                                                                                                 | Use in this study                                                                                                                                                                                                                                 |
|-------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| MIT–BIH<br>Normal Sinus<br>Rhythm DB<br>(NSRDB) | 18 real long-term ECG recordings from five men aged 26–45 and 13 women aged 20–50, with no significant arrhythmias; two stored channels named ECG1 and ECG2; 128 Hz; approximately 23–26 hours per record [6]. | The complete first stored ECG channel from every record is processed. Fifteen subjects fit the normal-only autoencoder and three different subjects validate it and determine the normal residual threshold.                                      |
| MIT–BIH<br>Arrhythmia DB<br>(MITDB)             | 48 approximately 30-minute, two-channel ambulatory ECG records from 47 subjects (25 men aged 32–89 and 22 women aged 23–89), sampled at 360 Hz, with expert beat and rhythm annotations [5, 2].                | Only the 46 records containing MLII are eligible. Their earlier 80% supplies supervised development and their final 20% supplies the temporal test. These records represent 45 subject clusters because records 201 and 202 belong to one person. |

Figure 3.1 gives the same information as a flow diagram. In short, NSRDB contributes only healthy windows for the frozen autoencoder, while MITDB contributes MLII beats for supervised development and the final temporal test.

### Dataset roles and evaluation boundaries



The temporal test uses later signal from the same MITDB patients; it is not an unseen-patient test.

**Figure 3.1:** Dataset roles and evaluation boundaries. The healthy NSRDB channel creates the frozen normal model. MITDB MLII creates the supervised feature table and is split chronologically into development and temporal-test regions.

**Why NSRDB is used for healthy training.** NSRDB contains real long-duration ECG from subjects reported to have no significant arrhythmias. This matches the autoencoder idea: the network sees normal structure without supervised arrhythmia labels. This is a nominally normal source cohort, not an annotation-verified guarantee that every window is normal. The preparation code does not exclude windows using NSRDB beat annotations. All 18 subjects and their complete recordings enter the cleaning process, giving 437.49 hours of source ECG. Complete recordings also expose the model to normal within-day variation instead of a short selected excerpt.


**Exactly one channel is used.** NSRDB headers call the two waveforms ECG1 and ECG2; they do not identify either as a standard anatomical lead. This study therefore selects the first stored channel consistently and describes it as the *first NSRDB ECG channel*. It is not renamed Lead II. The second channel never enters autoencoder fitting, validation, threshold estimation, or residual calculation.

The supervised branch uses MLII only. Modified Lead II is a bipolar torso lead with a negative RA-equivalent electrode on the right upper torso and a positive LL-equivalent electrode on the left lower torso. Figure 3.2 shows this polarity schematically; the separate

reference electrode is shown for the project's sensor connection. It does not represent a second ECG channel. PhysioNet describes the modified torso lead and its generally prominent normal QRS complexes [25, 5]. The loader selects MLII by its header name rather than assuming a channel number. MLII is stored first in 45 eligible records; record 114 stores V5 first and MLII second. Records 102 and 104 contain V5 and V2 but no MLII and are excluded. Consequently, the detector, classifier, temporal test, and installed replay all receive one MLII waveform.

The two model stages are one-channel systems, but the channel identities are not falsely treated as identical: the autoencoder learns healthy structure from the first NSRDB channel, while the supervised hierarchy operates on MLII. This source difference is a limitation because the electrode placement for the NSRDB channel is not documented.

### A MLII polarity and torso placement



Front view: patient's right is on the left of the drawing.

### B One channel per input; channel identities differ

#### Normal autoencoder: NSRDB

First stored channel: ECG1  
Anatomical placement is unspecified.  
Used to learn a normal ECG reference.

#### Supervised models: MITDB

Select MLII by its header name.  
Record 114: MLII is the second channel.  
Exclude 102 and 104: neither has MLII.

**Figure 3.2:** MLII polarity and the channel used by each model. Panel A is an anterior-view schematic of the modified Lead-II axis, with the project's reference connection shown separately; it is not an exact map of the database electrode coordinates [25]. Panel B distinguishes the NSRDB ECG1 normal reference from the MLII input used by the supervised models.

**Reference labels.** MITDB expert annotations are mapped to the six project outputs as shown in Table 3.2. AFib is determined from the rhythm annotation around a beat; the other targets mainly use beat symbols.

**Table 3.2:** Mapping from MITDB expert annotations to project outputs.

| Output | MITDB evidence                          | Interpretation                       |
|--------|-----------------------------------------|--------------------------------------|
| N      | N, e, j outside AFib/AFL                | normal-like/escape group             |
| PVC    | V, E                                    | ventricular premature or escape beat |
| PAC    | A, a, J, S                              | supraventricular premature beat      |
| LBBB   | L                                       | left bundle branch block beat        |
| RBBB   | R                                       | right bundle branch block beat       |
| AFib   | normal-like beat inside AFIB/AFL rhythm | atrial fibrillation/flutter context  |

The class codes are project groupings, not exact clinical diagnoses or the standard AAMI five-class scheme. N includes atrial and junctional escape beats; PVC includes ventricular escape beats; PAC includes nodal and supraventricular premature beats. AFib combines fibrillation and flutter rhythm annotations only when the beat symbol maps to N. Other supported beat types retain their morphology label during AFIB/AFL. Consequently, AFib recall here is neither pure AF sensitivity nor episode-level AF detection.

Unsupported symbols remain abnormal in binary evaluation and are excluded from six-class scoring. At runtime, a gated unsupported abnormality still receives one of the five available subtype labels. The MLII cache contains 104,755 accepted beats: 63,073 normal and 41,682 abnormal. Of these, 100,072 have a supported six-class label.

## 3.2 Data Cleaning

Data cleaning converts two different raw ECG sources into the same numerical input: one four-second signal with 512 samples. The cleaning rules are fixed before evaluation and are reused after installation. Figure 3.3 shows the full transformation for one healthy NSRDB segment and one MITDB MLII beat.

### 3.2.1 Cleaning Checklist

The cleaning pipeline is easier to check as a list. Each item has one purpose and one failure rule.

**Table 3.3:** Cleaning checklist applied before a signal can become a model input.

| Step                      | What is done                                                                                                                                                                           | Why it matters                                                                                                 |
|---------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------|
| Required channel          | NSRDB must provide ECG1; MITDB must provide MLII. Missing required signals stop processing.                                                                                            | Prevents accidental use of ECG2, V5, V2, or another non-matching lead.                                         |
| Sampling rate             | NSRDB remains at 128 Hz. MITDB MLII is resampled from 360 Hz to 128 Hz with polyphase anti-aliasing.                                                                                   | Makes every four-second model window exactly 512 samples.                                                      |
| Band-pass filter          | A Butterworth band-pass with prototype order four and cutoffs 0.5–50 Hz uses forward–backward filtering. For supervised inputs this is applied within the four-second window [26, 27]. | Reduces baseline movement and high-frequency noise without shifting the QRS complex.                           |
| Amplitude standardisation | Each window is converted to approximately zero mean and unit standard deviation.                                                                                                       | Reduces gain differences while preserving the relative P–QRS–T shape.                                          |
| Window quality check      | Corrected supervised windows must be finite and non-flat. Both neighbouring R peaks must be observed.                                                                                  | Rejects unusable inputs. These checks do not establish clinical signal quality or remove all motion artifacts. |

### 3.2.2 Window Construction

The autoencoder and supervised hierarchy build windows in different ways, but the final array shape is the same. The autoencoder receives consecutive, non-overlapping four-second windows from the first stored NSRDB channel. The supervised MLII windows are centred on a beat: 200 samples before the R peak and 312 samples after it. During offline experiments, expert annotations give the R-peak location for alignment and later scoring. During live use, detected R peaks replace those annotations.

For samples  $x_1, \dots, x_n$  inside one prepared model window, amplitude standardisation is

$$z_i = \frac{x_i - \mu}{\sigma + 10^{-8}}, \quad \mu = \frac{1}{n} \sum_{i=1}^n x_i. \quad (3.1)$$

The corrected supervised pipeline filters each extracted window and applies this operation locally. The original experiment instead standardised each complete MITDB record before extracting windows. That used test-region signal statistics during preparation and differed

from sensor chunks. The corrected results in this thesis replace that preprocessing. NSRDB remains the frozen reference checkpoint, with block filtering followed by per-window scaling.


### 3.2.3 Cleaning Yield

The healthy NSRDB recordings produce many candidate windows because complete records are used instead of short excerpts. Across the 18 complete NSRDB records, 393,744 candidate windows were formed. The quality checks kept 345,850 windows and rejected 47,894. These accepted windows were then split into 288,673 autoencoder fitting windows and 57,177 validation windows.

MITDB cleaning produces beat-centered MLII examples. The accepted MLII cache contains 104,755 beats: 63,073 normal and 41,682 abnormal. Unsupported abnormal symbols are still used for binary abnormality evaluation, but they are excluded from six-class subtype scoring.

### 3.2.4 Installed Use

The corrected hierarchy applies the same local window filtering, standardisation and feature order during fitting, offline scoring and installed use. Database replay first resamples the signal. Sensor acquisition supplies 128 Hz samples. Labels never enter a prediction. The live backend waits until there is enough signal around a detected R peak and enough RR context. It then feeds one prepared MLII window to the frozen system. Neither supervised model is retrained during prediction.



**Figure 3.3:** Cleaning real ECG for the two roles. The healthy example is the first channel of NSRDB record 16265. The unseen-use example is MLII from MITDB record 100. Both become one 512-sample input, but their documented channel identities remain different.

### 3.3 Data Preparation

Data preparation decides which cleaned examples are allowed to fit a model, which examples choose model settings, and which examples are held back for the final temporal result. Figure 3.4 shows this flow from raw source to frozen prediction.

#### 3.3.1 Prepared Objects

After cleaning, the study keeps two kinds of prepared object:

- **NSRDB window:** a healthy  $[1, 512]$  window from the first stored ECG channel. It has no arrhythmia label and is used only by the normal autoencoder.
- **MITDB beat:** an R-aligned MLII  $[1, 512]$  window with record identity, binary normal/abnormal truth, supported subtype truth when available, and RR timing context.

The autoencoder input tensor has shape  $[B, 1, 512]$ . The supervised detector and subtype

classifier do not receive raw samples alone. Each MITDB beat is converted into the same 242-value MLII feature vector:

$$\begin{aligned}
 242 \text{ values} &= 230 \text{ morphology, spectrum, statistics, and base-RR values} \\
 &+ 8 \text{ causal RR-history values} \\
 &+ 4 \text{ autoencoder residual values.}
 \end{aligned}$$

The subtype model is fitted only on the five supported abnormal classes: PVC, PAC, LBBB, RBBB, and AFib.

### 3.3.2 Splits

The autoencoder split is subject-disjoint. A seeded split assigns 15 complete NSRDB records to fitting and three different records (16420, 16483, and 16795) to validation. This produces 288,673 fitting windows and 57,177 validation windows. Validation selects the checkpoint and estimates the normal residual threshold. There is no separate healthy test because the autoencoder is frozen before it contributes residual features to MITDB evaluation.

The MITDB split is chronological inside every eligible record:

- the first 64% fits candidate supervised models;
- the next 16% selects the model structures and abnormal-probability threshold;
- a 16-beat gap is removed before the test boundary; and
- the final 20% evaluates the completed frozen hierarchy.


Percentages count accepted beats in chronological order, not exact elapsed time. The final run removes a 16-beat gap before both selection-validation and test boundaries and checks sample separation. The test had been observed during earlier development, making this a retrospective known-patient evaluation with possible optimism.

### 3.3.3 Exact Model Allocation

**Table 3.4:** Training, validation, and temporal-test data supplied to each model.

| Model                | Candidate training                                           | Validation                                             | Final fitting and testing                                                                         |
|----------------------|--------------------------------------------------------------|--------------------------------------------------------|---------------------------------------------------------------------------------------------------|
| Autoencoder          | 288,673 first-channel windows from 15 complete NSRDB records | 57,177 windows from 3 different complete NSRDB records | Frozen while every MLII test beat produces four residual features                                 |
| Abnormality detector | 66,285 MLII beats: 40,469 normal and 25,816 abnormal         | 16,029 beats: 9,664 normal and 6,365 abnormal          | Fitted on 83,050 development beats; tested on 20,969 temporal-test beats                          |
| Subtype classifier   | 22,844 supported abnormal MLII beats                         | 5,651 supported abnormal beats                         | Fitted on 28,758 supported abnormal beats; conditionally tested on 7,958 supported abnormal beats |

The supported six-output test contains 20,043 beats: 12,085 N, 1,527 PVC, 589 PAC, 1,481 LBBB, 1,383 RBBB, and 2,978 AFib. Another 926 unsupported abnormal beats remain in binary testing but not six-class scoring.



**Expert labels are stored separately from the 242 input features.**

Earlier MITDB beats develop the models; later beats evaluate their predictions.  
The first and last eligible beat are excluded when a neighbouring RR is unavailable.

**Figure 3.4:** Preparation of normal-reference windows and labelled MLII beats. Panel A develops the frozen autoencoder from subject-disjoint NSRDB fitting and validation data. Panel B uses local filtering and standardisation for each MITDB beat, then combines waveform, RR timing, and frozen-autoencoder residual features. Expert labels remain separate from the model input.

### 3.3.4 Feeding the Frozen System

During autoencoder fitting, Gaussian noise with standard deviation 0.03 is added to each input while the clean NSRDB window remains the target. During supervised fitting, record-and-class weights reduce domination by long records and common classes. During prediction, one MLII window and its causal RR history are converted into the same 242-value vector used in the temporal experiment. A normal decision stops. An abnormal decision reuses that vector for PVC, PAC, LBBB, RBBB, or AFib classification.


# 4 Theoretical Foundations

This chapter explains the theory behind the final ECG system. It follows the actual implementation: a one-lead MLII beat classifier that uses waveform shape, RR timing, autoencoder residuals, and two supervised Extra Trees ensembles. The chapter does not assume a second ECG lead or an offline record-summary input, because the deployed model does not use those inputs.

## 4.1 Why ECG evidence has more than one form

An ECG is a voltage signal over time. For arrhythmia classification, useful evidence appears at two related scales. The first is beat-scale morphology: what the P–QRS–T shape looks like around one R peak. The second is rhythm timing: when the beat occurs relative to neighbouring beats. The final model uses both because they answer different questions.

### Two ECG evidence scales used by the final model



The final feature vector joins these scales with four normal-reconstruction residuals.


**Figure 4.1:** Two evidence scales used by the final model. Morphology describes the local shape of a beat; RR timing describes rhythm context around the beat.

PVC, LBBB, and RBBB can strongly affect QRS shape, width, slopes, and polarity. PAC and AFib may be easier to detect when timing evidence is also present: premature beats create short RR intervals, pauses create long RR intervals, and AFib creates repeated irregularity. Published heartbeat-classification work also shows the value of combining waveform morphology with heartbeat-interval features [3, 7].

## 4.2 Feature extraction theory

The raw 512-sample MLII window contains more detail than the final tree models need. Small changes in R-peak alignment or noise can change many raw sample values. The system therefore converts the window into compact measurements that keep useful ECG evidence while reducing unnecessary sample-level detail.

### Transparent operations that create the final feature vector



**Figure 4.2:** Feature-extraction operations behind the final 242-value vector. The tree models learn from compact numerical evidence, not from raw ECG samples directly.

For a pooled region  $B_k$ , the pooled value is

$$z_k = \frac{1}{|B_k|} \sum_{t \in B_k} x_t. \quad (4.1)$$

Pooling keeps the average shape in a small region and makes the representation less sensitive to a one- or two-sample shift. The final system uses 64 whole-window pooled values and 80 central morphology pooled values.

The derivative

$$\Delta x_t = x_t - x_{t-1} \quad (4.2)$$

describes how quickly voltage changes. This matters because QRS abnormalities often change slopes and sharpness. The system keeps 60 pooled derivative values from the central beat region.

Frequency features describe how signal energy is spread from 0.5 to 50 Hz. For a frequency band  $b$ , the relative band feature is

$$s_b = \log \left( 1 + \frac{P_b}{P_{0.5-50}} \right), \quad (4.3)$$

where  $P_b$  is the power inside the band and  $P_{0.5-50}$  is total power in the ECG band used by the feature extractor. This band-limited view follows the general ECG idea of keeping physiological content while reducing drift and high-frequency interference [1, 26].

Eleven summary statistics add scale, range, energy, skewness, and kurtosis. Together, the waveform, derivative, spectral, and statistical values form 223 non-RR, non-autoencoder features.

### 4.3 RR timing theory

Let  $r_i$  be the time of the R peak of beat  $i$ . The RR interval is

$$RR_i = r_i - r_{i-1}. \quad (4.4)$$

RR values describe rhythm. A short interval can suggest a premature beat. A long interval can suggest a pause. Repeated interval changes support rhythm irregularity.

The final live system classifies a beat only after it has enough context to compute the same RR information used in the final experiment. It uses seven base RR values, including previous RR, next RR, local mean RR, interval ratios, local coefficient of variation, and heart rate. It then appends eight history-based values from already completed intervals: median, interquartile range, median absolute deviation, RMSSD, pNN50, relative range, trend, and latest interval change.

RMSSD is one example:

$$\text{RMSSD} = \sqrt{\frac{1}{K-1} \sum_{j=2}^K (RR_j - RR_{j-1})^2}. \quad (4.5)$$


The next-RR requirement is important. It creates a small beat-level delay during streaming, but it also makes the live RR feature vector match the experiment.

### 4.4 Normal-only reconstruction theory

An autoencoder contains an encoder  $f_\theta$  and decoder  $g_\phi$ . It is trained to reconstruct a normal ECG window  $x$ :

$$\hat{x} = g_\phi(f_\theta(x)), \quad \min_{\theta, \phi} \frac{1}{T} \sum_{t=1}^T (x_t - \hat{x}_t)^2. \quad (4.6)$$

This follows the representation-learning idea of autoencoders [14]. The final autoencoder is trained with a denoising objective: small Gaussian noise is added to the input while the target remains the clean normal window. Denoising encourages stable structure instead of exact noise copying [15].

**Normal-only reconstruction: residual is evidence, not diagnosis**

The supervised detector receives four residual summaries and decides together with morphology and RR evidence.

**Figure 4.3:** Normal-only reconstruction principle. A larger residual can indicate novelty, but it is evidence only; the final arrhythmia decision is made by the supervised hierarchy.

If the autoencoder has learned common normal structure, an unfamiliar or abnormal beat may be reconstructed less accurately. This is the one-class anomaly detection idea [16]. It is not a guarantee. A flexible decoder may reconstruct an abnormal waveform well, and noise may create a large residual for a normal waveform. Therefore, the final system uses autoencoder mismatch as four supervised input features, not as the final diagnosis rule.

Let  $e_t = x_t - \hat{x}_t$  be the residual. The four residual features are:

$$e_{\text{MAE}} = \frac{1}{T} \sum_t |e_t|, \quad (4.7)$$

$$e_{\text{MSE}} = \frac{1}{T} \sum_t e_t^2, \quad (4.8)$$

$$e_{\text{local}} = \frac{1}{|Q|} \sum_{t \in Q} |e_t|, \quad (4.9)$$


$$e_{\Delta} = \frac{1}{T-1} \sum_{t=2}^T |e_t - e_{t-1}|. \quad (4.10)$$

They measure whole-window mismatch, squared mismatch, central morphology mismatch, and residual sharpness. The final architecture appends these four values to the MLII waveform and RR features.

## 4.5 Hierarchical classification theory

The final problem is not solved with one direct classifier. The system first asks whether the beat is abnormal. Only when this binary gate is true does it ask which supported abnormal subtype is most likely.

### Hierarchical probability logic in the final model



**Figure 4.4:** Probability logic of the final hierarchy. The subtype model is conditional: it is used only after the abnormality gate is passed.

For beat  $i$ , the binary detector estimates  $p_i = P(\text{abnormal} | x_i)$ . The final threshold rule is

$$a_i = \mathbb{I}[p_i > 0.5708333]. \quad (4.11)$$

If  $a_i = 0$ , the output is N. If  $a_i = 1$ , the subtype model selects the highest probability among PVC, PAC, LBBS, RBBB, and AFib:

$$\hat{y}_i = \arg \max_{c \in \{1, \dots, 5\}} q_{ic}. \quad (4.12)$$

Here the maximisation already returns an index 1–5. Unsupported abnormal symbols contribute to binary testing and are excluded from subtype training and six-class scoring. At inference, every gated beat receives a supported label. There is no validated unknown-class rejection mechanism.

## 4.6 Why Extra Trees fits this feature design

The final 242-value vector mixes pooled waveform values, derivatives, spectral power, signal statistics, RR timing, and autoencoder residuals. Their relationship to class labels is non-linear. Extra Trees fits this setting by building many randomised decision trees and averaging their probability outputs [28].

At a decision-tree node with class probabilities  $p_c$ , entropy is

$$H = - \sum_c p_c \log_2 p_c. \quad (4.13)$$

A split is useful when it reduces the weighted child entropy. Extra Trees adds randomness by selecting random feature subsets and random split candidates. Averaging many such trees reduces the instability of a single tree while retaining non-linear interactions between features.

This theory also explains why Extra Trees is suitable for the final data scale. MITDB contributes many beats but only 45 subject clusters. A very high-capacity end-to-end neural classifier could learn patient identity instead of robust ECG rules [8]. The Extra Trees hierarchy is still high-capacity, but it uses transparent tabular evidence and is easier to audit against the final temporal protocol. This does not mean that tree ensembles are always better than deep learning.

## 4.7 Class and patient imbalance theory

ECG beat datasets are imbalanced in two ways. Some classes are rare, and some patients contribute many more beats than others. Without weighting, a model can mostly learn common normal beats or a few large records.

The final training code uses record-aware class weights. For class  $c$  in subject cluster  $r$ , the weight is proportional to

$$w_{cr} = \frac{1}{n_c} \left( \frac{n_c}{R_c n_{cr}} \right)^\alpha, \quad (4.14)$$

where  $n_c$  is total support for class  $c$ ,  $R_c$  is the number of subject clusters containing class  $c$ , and  $n_{cr}$  is the number of class- $c$  beats in record  $r$  (records 201 and 202 are combined). Weights are normalised to mean one and clipped to  $[0.05, 20]$ . The first factor supports class balancing; the second factor reduces domination by a single record. The exponent  $\alpha$  controls how strongly record balancing is applied and is selected in validation. Class-balanced learning is important when rare classes would otherwise have little influence [22].

## 4.8 Temporal patient adaptation theory

ECG morphology contains patient-specific information such as body geometry, electrode placement, baseline QRS shape, conduction pattern, and background rhythm. If earlier labelled ECG from a patient is available during development, the model can adapt to that patient's baseline. Later ECG from the same patient is then easier than ECG from a completely unseen patient.


This is the theory behind the final temporal-holdout design. It is useful for monitoring after a labelled calibration period, but it changes the claim. The final result is a known-patient temporal-continuation result. It should not be reported as unseen-patient generalisation unless a patient-disjoint experiment is run. The 16-beat boundary gap reduces direct overlap at the split boundary, but it does not remove shared patient identity.

## 4.9 Validation theory


Beat-level accuracy alone is not enough. If one patient contributes 2,000 beats, those beats are related. They are not equivalent to 2,000 independent patients. The final validation therefore reports beat metrics and also uses subject-clustered uncertainty.

### Why validation must respect subject clusters

#### A. Many beats, but fewer independent subjects



#### B. Bootstrap resamples subjects, not individual beats



This supports a known-patient temporal-continuation claim, not an unseen-patient claim.

**Figure 4.5:** Theoretical reason for subject-clustered validation. Clustered resampling prevents one long record from being treated as many independent people.

Balanced accuracy gives equal weight to class recalls, so normal beats cannot hide poor abnormal recall. Macro F1 gives each class equal importance, so a failure on a rare class is visible. Precision–recall style measures are also important under class imbalance [21]. For confidence intervals, the final experiment resamples complete subject clusters rather than individual beats, following the general diagnostic-statistics principle that clustered data need clustered analysis [23].

## 4.10 Theory-to-system alignment

**Table 4.1:** How the theory maps to the implemented final system.

| Theory concept             | Implemented in the final system                                                               |
|----------------------------|-----------------------------------------------------------------------------------------------|
| Beat-scale morphology      | 512-sample MLII window with pooled full-window, central, and derivative features              |
| Rhythm timing              | 15 RR features; live scoring waits until next-RR context is available                         |
| Normal-only reconstruction | frozen NSRDB autoencoder supplies four residual features                                      |
| Hierarchical probability   | 240-tree binary gate first; 260-tree subtype model only for gated abnormal beats              |
| Class/patient imbalance    | record-aware class weighting with validation-selected record-power parameter                  |
| Temporal adaptation        | earlier MLII timeline trains/selects the model; final 20% of accepted beats is tested         |
| Subject-aware uncertainty  | subject-clustered bootstrap over 45 clusters, with records 201 and 202 treated as one cluster |


# 5 Final Model Architecture

This chapter explains the final system saved and used by the application. The installed model is a one-lead MLII beat classifier. It does not use two ECG leads, does not retrain during live use, and does not make an arrhythmia decision directly from the autoencoder threshold. The final decision is made by two supervised Extra Trees models. Both receive the same 242-value feature vector.

## 5.1 Architecture in one sentence

For each detected MLII heartbeat, the system takes a 512-sample window aligned to the R peak, builds waveform, rhythm, and normal-reconstruction features, sends those features to a binary abnormality gate, and sends only gated abnormal beats to a five-class subtype classifier.

### Final ECG system: from one MLII beat to one decision



The autoencoder supplies four residual features. It does not make the final class decision alone.

**Figure 5.1:** Exact runtime order of the final ECG system. The autoencoder supplies features; the supervised tree models make the final class decision.

## 5.2 Input signal and beat window


The deployed system accepts one ECG channel. For database replay and final testing that channel is MIT-BIH MLII, selected by signal name. Records 102 and 104 are excluded because they do not contain MLII. For sensor sessions, the API requires the user to declare

`lead_name='MLII'` and rejects other declared leads. This rule prevents the model from silently scoring a lead that does not match the supervised training signal.

The classifier is beat based, not arbitrary-window based. The live backend collects MLII samples, detects R peaks, and waits until enough rhythm context is available. For a scorable beat, it extracts a fixed 512-sample window at 128 Hz:

- 200 samples before the R peak;
- 312 samples after the R peak;
- total duration of 4.0 seconds.

This same window geometry is used in the final temporal-holdout experiment and in the installed session pipeline.



**Figure 5.2:** R-peak-aligned beat window used by the final model. The R peak is at sample index 200, giving the model pre-QRS context and post-QRS recovery information.


### 5.3 Three trained components

The final system uses three trained components, stored as an autoencoder checkpoint and a separate hierarchy bundle:

1. a frozen normal-only autoencoder;
2. a binary Extra Trees abnormality detector;
3. a five-class Extra Trees abnormal-subtype classifier.

The autoencoder is trained first and then frozen. The two supervised tree models are then fitted on MLII features from the temporal development portion of MIT-BIH Arrhythmia Database records.

### A Offline training produces three fixed components



**Figure 5.3:** Training and use of the saved three-component model. Offline development produces the frozen autoencoder, binary gate, subtype classifier, and validation-selected threshold. The same fixed components support retrospective temporal evaluation and session prediction; session use does not update its parameters.

## 5.4 Normal-only autoencoder

### 5.4.1 Purpose

The autoencoder is a normal-reconstruction branch. It learns how normal ECG windows can be reconstructed. The final system then uses reconstruction error as extra evidence. The autoencoder does not decide “normal” or “arrhythmia” by itself in the final classifier.

For an input window  $x \in \mathbb{R}^{512}$ , the autoencoder output is


$$\hat{x} = g_{\phi}(f_{\theta}(x)), \quad (5.1)$$

where  $f_{\theta}$  is the encoder and  $g_{\phi}$  is the decoder. Training uses a denoising objective: Gaussian noise with standard deviation 0.03 is added to the input, but the target remains the clean standardised window. The training loss is

$$\mathcal{L}_{AE} = \frac{1}{512} \sum_{t=1}^{512} (x_t - \hat{x}_t)^2. \quad (5.2)$$

This objective encourages stable normal morphology instead of copying small noise patterns [15].

### 5.4.2 Network structure



Dimensions are channels x time samples; the batch dimension is omitted.

Dashed arrows carry encoder features to matching decoder lengths.

Each up block resizes to its skip length; final interpolation restores 512 samples.

**Figure 5.4:** Autoencoder architecture with tensor dimensions verified by a forward pass through the 2,680,641-parameter implementation. Read down the encoder, through the bottleneck, and up the decoder. Dashed connections concatenate encoder features scaled by 0.5; decoder outputs match their skip lengths before the final interpolation restores the 512-sample waveform.


Each residual encoder block contains two kernel-seven convolutions, Group Normalisation, LeakyReLU, dropout, and a residual projection. Four kernel-three bottleneck blocks use dilation 1, 2, 4, and 8. Dilation increases the receptive field without deleting time samples. The decoder uses transposed convolutions and skip connections scaled by 0.5. Each up block resizes its output to match the corresponding encoder skip. The first up block therefore returns

length 32, the fourth returns length 256, and final linear interpolation restores length 512. The figure reports these actual tensor dimensions. The final layer is linear because the standardised ECG signal contains positive and negative values.

### 5.4.3 Normal-reference data and checkpoint

The autoencoder checkpoint was developed from complete recordings of all 18 subjects in the MIT-BIH Normal Sinus Rhythm Database. Fifteen subjects supply 288,673 fitting windows and three different subjects supply 57,177 validation windows. The selected NSRDB header channel is ECG1. The database does not document this channel as anatomical Lead II, so the report treats it as a normal ECG reference channel rather than renaming it MLII [6].

The final checkpoint uses Adam, learning rate 0.001, batch size 256, and two epochs. Validation MSE changed from 0.061195 after epoch 1 to the best saved value, 0.061193, after epoch 2. The autoencoder artifact also stores a normal-residual threshold of 0.158893, but the final hierarchy does not use that threshold as the final arrhythmia rule.



**Figure 5.5:** Autoencoder fitting and subject-disjoint validation loss. The checkpoint with the lowest validation loss is retained before the supervised MLII stages are trained.

#### 5.4.4 Four autoencoder residual features

Let  $e_t = x_t - \hat{x}_t$  be the reconstruction residual. The final feature system uses four scalar residual summaries:


$$f_1 = \frac{1}{512} \sum_t |e_t| \quad \text{whole-window MAE,} \quad (5.3)$$

$$f_2 = \frac{1}{512} \sum_t e_t^2 \quad \text{whole-window MSE,} \quad (5.4)$$

$$f_3 = \frac{1}{240} \sum_{t=104}^{343} |e_t| \quad \text{central morphology MAE,} \quad (5.5)$$

$$f_4 = \frac{1}{511} \sum_{t=2}^{512} |e_t - e_{t-1}| \quad \text{residual-change MAE.} \quad (5.6)$$

The central feature focuses on the QRS and nearby morphology. The residual-change feature becomes large when the reconstruction misses sharp slopes. These four values are appended to the transparent MLII and RR features.




**Figure 5.6:** Real healthy and abnormal ECG examples through the frozen autoencoder. The residual between input and reconstruction is converted into four numerical features for the supervised detector.

## 5.5 Final 242-value feature vector

The saved final bundle declares 242 total features, 15 RR features, required autoencoder features, and four autoencoder residual values. Therefore every supervised prediction uses:

$$\begin{aligned}
 &223 \text{ waveform/spectrum/statistics values} \\
 &+ 15 \text{ RR timing values} \\
 &+ 4 \text{ autoencoder residual values} = 242.
 \end{aligned}
 \tag{5.7}$$



**Figure 5.7:** Feature-vector layout used by both supervised models. The same 242-value vector is passed first to the abnormality detector and, if required, to the subtype classifier.

**Table 5.1:** Organised feature groups in the final architecture.

| Group                   | Values | What it tells the model                                                                                                                   |
|-------------------------|--------|-------------------------------------------------------------------------------------------------------------------------------------------|
| Whole-window mean bins  | 64     | broad four-second MLII waveform shape                                                                                                     |
| Central morphology bins | 80     | P–QRS–T shape around the R peak                                                                                                           |
| Central derivative bins | 60     | sharp slopes and rapid QRS changes                                                                                                        |
| Frequency-band powers   | 8      | relative energy from 0.5 to 50 Hz                                                                                                         |
| Signal statistics       | 11     | scale, range, energy, skewness, and kurtosis                                                                                              |
| RR timing context       | 15     | previous/next interval, local rate, variability, median, interquartile range, MAD, RMSSD, pNN50, range ratio, trend, and latest RR change |
| Autoencoder residuals   | 4      | how strongly the beat disagrees with the learned normal-reconstruction pattern                                                            |

The final experiment constructs the same 242 values in two steps: first 230 values are built from 223 MLII waveform/statistical values plus seven base RR values; then eight causal RR-history values and four autoencoder residual values are appended. During live/session prediction, the backend builds the same final vector directly from 223 MLII waveform/statistical values, 15 RR values, and four autoencoder residual values.

## 5.6 Binary abnormality detector

The first supervised model is the abnormality detector. It receives all 242 features and estimates  $p_i = P(\text{abnormal} \mid x_i)$  for beat  $i$ . The selected detector is an Extra Trees classifier with:

- 240 trees;
- entropy split criterion;
- fully grown trees;
- minimum leaf size of one sample;
- square-root feature sampling at each split;
- record/class weighting with record-power  $\alpha = 0.35$ .

The decision threshold was selected on the 64–80% validation region of each development record:

$$a_i = \mathbb{I}[p_i > 0.5708333]. \quad (5.8)$$

If  $a_i = 0$ , the final output is normal class N. No subtype classifier is used for that beat.

For tree  $b$ , let  $p_b(x_i)$  be the abnormal class proportion in the reached leaf. The ensemble probability is

$$p_i = \frac{1}{240} \sum_{b=1}^{240} p_b(x_i). \quad (5.9)$$

Averaging many randomly split trees reduces the instability of one decision tree while still modelling non-linear links among ECG shape, RR timing, and autoencoder residuals.

## 5.7 Abnormal-subtype classifier

The second supervised model is used only after the binary gate marks a beat as abnormal. It is fitted only on supported abnormal classes 1–5, not on N and not on unsupported MIT-BIH symbols. The selected subtype model is an Extra Trees classifier with:

- 260 trees;
- entropy split criterion;
- fully grown trees;
- minimum leaf size of one sample;
- one-third feature sampling at each split;
- record/class weighting with record-power  $\alpha = 0.25$ .

For an abnormal beat, the subtype classifier returns probabilities over the five supported abnormal outputs: PVC, PAC, LBBB, RBBB, and AFib. The final subtype is

$$\hat{y}_i = \arg \max_{c \in \{1, \dots, 5\}} q_{ic}. \quad (5.10)$$

The final saved bundle sets `reject_low_confidence=False`. Therefore the evaluated and installed final pipeline do not add a post-hoc “Unclassified abnormal” rejection step; every gated abnormal beat receives the subtype with the largest subtype probability.

The displayed subtype confidence is conditional on the abnormal branch. It is an ensemble score, not a calibrated probability of disease. Unsupported abnormalities can receive an incorrect supported subtype.

## 5.8 Final decision logic

The complete prediction rule is:

1. Validate that the input is a 512-sample MLII beat window and that 15 RR features are available.
2. Reconstruct the same window with the frozen autoencoder.
3. Build the 242-value vector from MLII shape, RR timing, and four autoencoder residual summaries.
4. Estimate  $p_i = P(\text{abnormal} \mid x_i)$  using the 240-tree binary detector.
5. If  $p_i \leq 0.5708333$ , output N with confidence  $1 - p_i$ .
6. If  $p_i > 0.5708333$ , run the 260-tree subtype model and output the highest-probability supported subtype.

This is the same hierarchy used in the final temporal-holdout evaluation. The model does not use test labels while predicting. The labels are used only after prediction to compute accuracy, F1, confusion matrices, and subject-clustered uncertainty.

## 5.9 Why this architecture was selected

**Table 5.2:** Reason for each major final-architecture decision.

| Decision                        | Reason                                                                                                                                |
|---------------------------------|---------------------------------------------------------------------------------------------------------------------------------------|
| MLII-only supervised input      | matches the signal available in 46 eligible MIT-BIH records and avoids hidden two-lead assumptions                                    |
| R-peak alignment                | puts comparable cardiac events in comparable positions inside every 512-sample window                                                 |
| Waveform and derivative bins    | preserve beat morphology, including sharp QRS changes                                                                                 |
| Frequency powers and statistics | add compact information about signal energy, scale, and distribution                                                                  |
| Fifteen RR features             | represent local rate, premature beats, rhythm irregularity, and recent interval trend using information reproducible during streaming |
| Normal-only autoencoder         | supplies a learned measure of disagreement with normal ECG structure without requiring arrhythmia labels                              |
| Two-stage hierarchy             | allows unsupported abnormal beats to count in binary testing without forcing them into one of the five supported subtype labels       |
| Extra Trees models              | handle mixed feature scales and non-linear interactions with reproducible tabular-model training                                      |
| Temporal split with gap         | evaluates later signal from each record while removing direct boundary overlap between development and test beats                     |

The architecture is therefore best described as a known-patient temporal continuation system. Its final performance numbers describe later MLII signal from patients whose earlier MLII signal was available during model development. They should not be presented as unseen-patient generalisation without a separate patient-disjoint evaluation.

# 6 Experimental Method

The final experiment measures one specific use case: the system has already seen earlier MLII ECG from a patient, and it must classify later ECG from that same patient. This is a known-patient temporal holdout. It is not a new-patient external validation and it is not a shuffled random beat split.

## 6.1 Method overview

Figure 6.1 shows how accepted beats are allocated to model selection, final refitting, and testing. The normal autoencoder was already frozen from NSRDB. The supervised experiment then used only MITDB MLII beats. Candidate models were selected with earlier temporal regions, refitted before the test boundary, and evaluated only on the final 20% of accepted beats.

### Chronological split within each MITDB record


Percentages refer to accepted-beat order, not elapsed recording time.

#### A Select model settings



Validation chooses the model candidates and the abnormality threshold.

#### B Refit, freeze, then evaluate



The model and threshold remain fixed while scoring the final test segment.

Hatched markers: 16 beats excluded immediately before each marked boundary. Markers are enlarged for visibility. The 64% gap is restored during final refitting.

Interpretation: later beats from known patients; the test had been seen in earlier project development, so this is a retrospective evaluation.

**Figure 6.1:** Chronological selection, refitting, and testing by accepted-beat order. Selection excludes 16 beats before both boundaries; refitting retains only the exclusion before the test. The test region was observed during earlier development.

## 6.2 Research objective and success conditions

The final experiment asks:

After the model has access to the earlier ECG of each patient, how accurately does it classify that patient's later ECG?

The primary outcome is whole-system six-class accuracy on the final 20% temporal holdout. The planned success conditions were:

1. pooled whole-system accuracy at least 90%;
2. lower bound of the equal-subject 95% confidence interval above 90%; and
3. total confidence-interval width no greater than three percentage points.

Per-class recall and macro F1 were also checked so that a high overall score could not hide failure on a smaller arrhythmia class.

## 6.3 Input data and feature construction

The method uses one ECG channel at every stage, but the two data sources have different roles. NSRDB supplies healthy windows for the frozen normal autoencoder. MITDB supplies MLII beat windows, binary labels, six-class labels, and the final temporal test. No NSRDB beat is scored in the final test.

**Table 6.1:** Exact evidence supplied to the final supervised models.

| Evidence block               | Count      | Role                                                                                                                                                                                           |
|------------------------------|------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| MLII morphology and spectrum | 223 values | Describes the beat-centred 512-sample MLII waveform. This is the direct ECG evidence for both supervised stages.                                                                               |
| RR timing context            | 15 values  | Describes previous RR, next RR, local timing, and recent rhythm history. In live use the beat is scored only after this next-RR context is available; labels are never used during prediction. |
| Frozen autoencoder residuals | 4 values   | Summarises how the frozen NSRDB-trained normal autoencoder reconstructs the same MLII beat.                                                                                                    |
| Total per MITDB beat         | 242 values | The binary detector and subtype classifier receive the same one-lead feature vector.                                                                                                           |

This feature construction matches the saved final model bundle: `lead_count` is one, `required_lead` is MLII, `rr_feature_count` is 15, and `ae_feature_count` is four.

The method therefore does not use a second MITDB lead during training, validation, temporal testing, or deployment.


## 6.4 Temporal split

The eligible supervised records are the 46 MITDB records that contain MLII. Records 102 and 104 were excluded before fitting because their available leads are V5 and V2, not MLII. Records 201 and 202 are treated as one subject cluster for uncertainty analysis because they belong to the same person.

Every eligible record is kept in chronological order. Percentages refer to accepted-beat counts, not exact elapsed recording time. With  $n_r$  accepted beats, the cuts are  $\lfloor 0.64n_r \rfloor$  and  $\lfloor 0.80n_r \rfloor$ :

1. the first 64% is the candidate-training region;
2. a 16-beat gap is removed before validation during candidate selection;
3. the 64–80% region is the validation region;
4. a 16-beat gap is removed before the final temporal test; and
5. the final 20% is the fixed temporal test.

For final fitting, the selected binary and subtype models are refitted on the earlier 80% development region after the 16-beat boundary gap before the test is removed. The test segment itself remains intact.



**Figure 6.2:** Temporal split audit for every eligible MITDB MLII record. Blue beats are available for final fitting, yellow beats are the boundary gap before the test, and red beats are the held-out final 20% temporal test.


The final held-out test contains 20,969 accepted beats. These are later beats from known patients, not beats from unseen patients. The method prevents signal overlap across the test boundary, but it intentionally keeps patient overlap.

## 6.5 Model selection and final fitting

Model selection uses only the earlier temporal validation region. In this rerun, the final 20% test region does not choose the candidate, threshold, or hyperparameters.

Three binary Extra Trees candidates were fitted on the earlier candidate region and ranked by validation balanced accuracy, then by F1. Three subtype Extra Trees candidates were fitted only on supported abnormal candidate beats and ranked by validation macro F1, then by accuracy. Figure 6.3 shows the validation scores that selected the final pair.

### Validation-region model selection




**Figure 6.3:** Model selection on the 64–80% validation region. The binary detector is selected by balanced accuracy and F1; the subtype classifier is selected by macro F1 and accuracy.

The selected binary detector was the fully grown 240-tree candidate with square-root feature sampling, minimum leaf size one, and record-weighting power 0.35. Its validation-selected abnormal threshold was 0.5708333. The selected subtype classifier was the fully grown 260-tree candidate with one-third feature sampling, minimum leaf size one, and record-weighting power 0.25. After selection, the binary model was refitted on all development beats, whereas the subtype model was refitted only on supported abnormal development beats.

## 6.6 Evaluation order

The deployed decision is hierarchical. First, the binary detector estimates whether a beat is abnormal. If the abnormal probability is not above the fixed threshold, the final output is N. If the beat passes the gate, the subtype classifier chooses PVC, PAC, LBBB, RBBB, or AFib.

### Three distinct evaluation populations



These groups overlap. Unsupported abnormal labels enter only binary evaluation.

**Figure 6.4:** Populations used for the three evaluation questions. Conditional subtype scoring uses true supported abnormal beats independently of the gate. Complete-system scoring uses the gate and subtype model together. The populations overlap; unsupported abnormal labels enter only the binary evaluation.

The temporal evaluation reports three supports:

1. binary abnormality detection on all 20,969 held-out beats, including unsupported abnormal MITDB symbols;
2. conditional subtype classification on 7,958 true supported abnormal beats; and
3. whole-system N/PVC/PAC/LBBB/RBBB/AFib classification on 20,043 supported held-out beats.

The 926 unsupported abnormal beats are useful for testing the binary normal/abnormal gate, but they are not scored in the six-class whole-system confusion matrix.

## 6.7 Metrics

For a binary outcome,

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}, \quad (6.1)$$

$$\text{Sensitivity} = \frac{TP}{TP + FN}, \quad (6.2)$$

$$\text{Specificity} = \frac{TN}{TN + FP}, \quad (6.3)$$

$$\text{Balanced accuracy} = \frac{\text{Sensitivity} + \text{Specificity}}{2}. \quad (6.4)$$

For class  $c$ ,

$$\text{Precision}_c = \frac{TP_c}{TP_c + FP_c}, \quad (6.5)$$

$$\text{Recall}_c = \frac{TP_c}{TP_c + FN_c}, \quad (6.6)$$

$$F1_c = \frac{2 \text{Precision}_c \text{Recall}_c}{\text{Precision}_c + \text{Recall}_c}. \quad (6.7)$$

Macro F1 is the unweighted mean of class F1 values. It gives PAC the same weight as N even though N contains many more beats.

Two accuracy summaries are reported. Pooled beat accuracy gives every beat equal weight. Subject-macro accuracy first calculates accuracy inside each of the 45 eligible subject clusters and then gives every subject equal weight:

$$A_{macro} = \frac{1}{45} \sum_{s=1}^{45} A_s. \quad (6.8)$$

The subject-macro value is primary for uncertainty because a long record should not behave like many independent patients.

## 6.8 Subject-clustered confidence interval

The primary 95% confidence interval uses a non-parametric percentile bootstrap:

1. keep every beat from one subject together;
2. treat records 201 and 202 as one subject cluster;
3. sample 45 subject clusters with replacement;
4. calculate the equal-subject mean accuracy;
5. repeat 10,000 times using seed 20260801; and
6. take the 2.5th and 97.5th percentiles.

These intervals describe variation across the evaluated subject clusters with the fitted model held fixed. They do not include retraining variability, uncertainty from earlier test exposure, or a new population shift. Resampling individual beats would create an artificially narrow interval because adjacent beats from the same recording are correlated. The clustered bootstrap keeps that dependence [23].

## 6.9 Tests against the 90% target

The confidence interval is the primary statistical evidence. Two non-parametric tests provide sensitivity analyses. The exact one-sided sign test asks whether more than half of

subjects exceed 90%. The one-sided Wilcoxon signed-rank test asks whether subject accuracies are shifted above 90%, under the additional symmetry assumption of a signed-rank location test. Accuracy is bounded and has ceiling effects, so the Wilcoxon analysis is secondary. The sign test concerns the fraction of subjects above the target, not pooled beat accuracy. These tests are interpreted together with the confidence interval, interval width, per-class recall, macro F1, and subject accuracy distribution.

## 6.10 Reproducibility and leakage controls

The final experiment saves the temporal boundary audit for every record, the candidate scores, selected models, probability threshold, final confusion matrix, per-class counts, per-subject accuracies, bootstrap seed, and headline metrics. The rerun used seed 20260801 and did not use the final 20% segment for candidate selection, threshold selection, or final refitting.

The final 20% segment had been observed during earlier system development, so the evaluation is a retrospective rerun on previously observed data. This prior exposure can influence design choices and cause optimism even when the present selection code excludes test rows. It is not a confirmatory first-use test, and the bootstrap does not correct development-selection bias. Patient overlap is also intentional: the test is later ECG from known patients. Therefore, the method supports a patient-adapted temporal-performance claim, not a patient-independent generalisation claim.

# 7 Results and Statistical Validation

This chapter reports the corrected retrospective temporal experiment in `experiments/corrected_temporal_holdout.json`. The rerun replaces record-wide supervised normalisation with local window preparation, removes unobserved neighbouring-RR substitutions, and saves beat-level predictions. The original August artifact is retained for provenance. No new candidate family or threshold was selected from the temporal-test results.


## 7.1 Evaluation population and supports

The final dataset contains 104,755 accepted beats from 46 MITDB MLII records, representing 45 subject clusters. Records 102 and 104 lack MLII and are excluded. Records 201 and 202 form one subject cluster. The first and last eligible beat in each record are excluded because a preceding or following accepted R peak is unavailable. Compared with the original cache, this removes 92 beats before splitting. Percentages refer to accepted-beat order.

### One temporal test set; three evaluation questions

#### 20,969 later beats from 45 patient clusters

12,085 normal + 7,958 supported abnormal + 926 other abnormal



Supported arrhythmias: PVC, PAC, LBBB, RBBB and AFib.


The three groups overlap; each answers a separate evaluation question.

**Figure 7.1:** Evaluation groups for the corrected retrospective test. Conditional subtype scoring bypasses the binary gate; complete-system scoring applies the hierarchy. Counts overlap, and records 201 and 202 share one patient cluster.

**Table 7.1:** Distinct denominators for the three evaluation questions.

| Evaluation          | Support and interpretation                                                                      |
|---------------------|-------------------------------------------------------------------------------------------------|
| Binary detection    | 20,969 beats: 12,085 N and 8,884 abnormal, including 926 unsupported abnormal beats.            |
| Conditional subtype | 7,958 true supported abnormal beats, scored by the subtype model regardless of the binary gate. |
| Whole system        | 20,043 supported beats, scored after the gate: N/PVC/PAC/LBBB/RBBB/AFib.                        |

### What data was used for the final temporal test

**Figure 7.2:** Reference-class composition of the corrected temporal test.

## 7.2 Model selection

Candidate selection uses the earlier 64–80% region after the boundary exclusions. The chosen binary model has 240 fully grown Extra Trees, square-root feature sampling, minimum leaf size one, entropy criterion, and subject/class weighting power 0.35. Its validation balanced accuracy is 98.73%. The chosen subtype model has 260 trees, feature sampling 0.33, minimum leaf size one, entropy criterion, and weighting power 0.25. Its validation macro F1 is 0.9811.


The binary threshold is 0.5708333333. The selected models are refitted on 83,050 development beats and 28,758 supported abnormal development beats, respectively. The final test does not change this threshold.

## 7.3 Headline temporal results

**Table 7.2:** Corrected temporal results. Confidence intervals estimate equal-subject accuracy, not pooled beat accuracy.

| Outcome               | Pooled | Subject mean | 95% CI       | Width    |
|-----------------------|--------|--------------|--------------|----------|
| Binary accuracy       | 98.81% | 98.84%       | 98.01–99.45% | 1.44 pp  |
| Conditional subtype   | 97.63% | 94.77%       | 89.03–99.08% | 10.05 pp |
| Whole-system accuracy | 98.07% | 98.03%       | 96.29–99.25% | 2.95 pp  |

Binary balanced accuracy is 98.69%, sensitivity is 97.91%, specificity is 99.47%, precision is 99.27%, and F1 is 0.9858. Conditional subtype macro F1 is 0.9574; whole-system macro F1 is 0.9524.




**Figure 7.3:** Pooled temporal scores after the preprocessing correction.

The three aggregate checks pass: pooled accuracy exceeds 90%, the equal-subject lower confidence bound exceeds 90%, and the interval width is at most three percentage points. These are descriptive outcomes of the retrospective study. Prior test exposure prevents interpreting them as an unbiased confirmatory success.


## 7.4 Uncertainty and subject variation

The bootstrap resamples 45 complete subject clusters with equal weight, 10,000 times using seed 20260801. It holds the fitted model fixed. For subtype accuracy, subjects without supported abnormal test beats have no subtype score and are omitted from that replicate's mean.



**Figure 7.4:** Equal-subject accuracy and percentile bootstrap intervals. The conditional subtype interval remains wide even though the whole-system width target is met.

Forty-three of 45 subject clusters exceed 90% whole-system accuracy. The exact one-sided sign test gives  $p = 2.94 \times 10^{-11}$  against the null that no more than half of subjects exceed the target. The secondary Wilcoxon statistic is 987, with  $p = 4.24 \times 10^{-8}$ . Its location-shift interpretation requires additional symmetry assumptions. Neither test corrects selection optimism or establishes new-patient performance.



**Figure 7.5:** Corrected whole-system accuracy across 45 known-patient clusters.


The below-target clusters are 207 (67.92%) and 222 (84.85%). Cluster 207 has high binary accuracy but poorer subtype classification. Cluster 222 has binary accuracy 85.45%, so failures in the gate also matter. This variation is not visible from pooled accuracy alone.

## 7.5 Errors by stage and class

The gate detects 8,698 of 8,884 abnormal beats and correctly rejects 12,021 of 12,085 N beats. Thus it produces 186 missed abnormal beats and 64 false alarms. The independent subtype model correctly names 7,769 of 7,958 supported abnormal beats. Gate errors can still send correctly classifiable abnormalities to N.


**Table 7.3:** Whole-system performance by supported abnormal output group.

| Group | True  | Correct | Precision | Recall | F1     |
|-------|-------|---------|-----------|--------|--------|
| PVC   | 1,527 | 1,499   | 92.36%    | 98.17% | 0.9517 |
| PAC   | 589   | 416     | 98.58%    | 70.63% | 0.8229 |
| LBBB  | 1,481 | 1,476   | 96.72%    | 99.66% | 0.9817 |
| RBBB  | 1,383 | 1,365   | 99.78%    | 98.70% | 0.9924 |
| AFib  | 2,978 | 2,879   | 98.26%    | 96.68% | 0.9746 |



**Figure 7.6:** True and correctly classified supported abnormal beats. PAC remains below the 90% recall target.

Of 589 PAC-group beats, 549 pass the binary gate, 441 receive the correct subtype when the gate is ignored, and 416 are correct end to end. PAC errors therefore arise in both stages. AFib denotes the project's AFIB/AFL rhythm grouping on normal-like beat symbols; the 96.68% recall is not an episode-level or AF-only sensitivity estimate.



**Figure 7.7:** Whole-system confusion matrix for 20,043 supported temporal beats.

The largest abnormal-subtype confusion is 113 PAC beats assigned to PVC. The gate sends 91 AFib-group and 40 PAC-group beats to N. Among reference N beats, 34 become LBBB and 26 become AFib. Unsupported abnormal beats are excluded from this matrix, but a deployed gated beat still receives a supported subtype because the model has no unknown-class rejection rule.

## 7.6 Limited external INCART check

After fitting the corrected hierarchy, its frozen threshold was applied to complete INCART records I01 and I02. These were the two lead-II recordings already cached locally at the start of the audit. Selection was by local availability, not measured performance. This is a convenience sample. INCART contains 75 half-hour records from 32 source patients, with 12 leads sampled at 257 Hz [29]. I01 and I02 both carry the header identifier *patient 1*: the two records represent one independent source patient.

Standard lead II is a practical external counterpart to MITDB MLII, not an identical lead. This check combines source, acquisition and lead differences. Scoring uses annotated beats, not live R-peak detections.

**Table 7.4:** Corrected model on complete INCART I01 and I02 recordings.

| Quantity                  | Value                 |
|---------------------------|-----------------------|
| Records / source patients | 2 / 1                 |
| Accepted annotated beats  | 5,416                 |
| Reference N / PVC         | 4,845 / 571           |
| Fixed abnormal threshold  | 0.5708333333          |
| Binary accuracy           | 94.11%                |
| Binary balanced accuracy  | 96.55%                |
| Sensitivity / specificity | 99.65% / 93.46%       |
| Binary precision / F1     | 64.22% / 0.7811       |
| AUROC / average precision | 0.9914 / 0.9152       |
| TP / TN / FP / FN         | 569 / 4,528 / 317 / 2 |
| Whole-system accuracy     | 92.85%                |
| End-to-end PVC recall     | 87.74%                |

Only N and PVC have reference support. The binary model detects 569 of 571 PVC beats but produces 317 false alarms among 4,845 N beats. End to end, 501 PVC beats receive the correct label. Binary accuracies by record are 94.47% for I01 and 93.74% for I02. High binary sensitivity therefore coexists with lower precision and subtype errors.

The stored six-label macro F1 is 0.3039 because it includes four zero-support classes with F1 set to zero. The mean over the two reference-supported classes is 0.9116. These summaries are not directly comparable to a six-class test with support in every class. PAC, LBBB, RBBB and AFib sensitivity cannot be estimated from this subset. A patient-cluster population confidence interval is not reported because there is only one source patient.

The August check used only 120 I01 beats and reported 98.33% binary accuracy. It remains historical evidence and is superseded here by the complete-record check of the corrected model. No exploratory threshold optimum from external labels is installed or reported as validation performance. Both external records have now been inspected and must not be reused as an untouched test in a future study.

## 7.7 Answers to the research questions

**RQ1** The frozen NSRDB autoencoder supplies four residual features. No controlled ablation establishes its independent benefit.

**RQ2** Binary, conditional-subtype and whole-system accuracies are 98.81%, 97.63% and 98.07%, respectively, on their stated supports.

**RQ3** Equal-subject whole-system accuracy is 98.03%, with a 96.29–99.25% interval and two below-target subject clusters.

**RQ4** The three aggregate accuracy/precision checks pass, but PAC recall does not. The

INCART convenience subset represents one source patient and only two reference classes, so broad external validation remains incomplete.

# 8 Discussion and Conclusion

---

## 8.1 Main findings

The corrected one-lead hierarchy achieves 98.07% pooled six-class accuracy on later accepted MITDB MLII beats from known patients. Equal-subject accuracy is 98.03%, with a subject-clustered 95% interval of 96.29–99.25%. The 2.95-point interval width meets the three-point precision target. However, PAC recall is 70.63%, and two subject clusters remain below 90% accuracy.

A limited check of complete INCART I01 and I02 records gives 94.11% binary accuracy and 92.85% whole-system accuracy. Both records belong to the same source patient. Binary precision falls to 64.22% because 317 normal beats become false alarms. These results support further external testing, not a clinical generalisation claim.

## 8.2 Achievement of the objectives

- The prototype integrates a frozen NSRDB autoencoder and two supervised MLII ensembles, with the same local model-window preparation at training and use.
- Pooled whole-system accuracy, the equal-subject lower confidence bound, and the interval-width target are met in the corrected retrospective run.
- The requirement that every supported abnormal group exceed 90% recall is not met because of PAC.
- Generalisation across independent external patients, all supported subtypes, and real AD8232 recordings is not established.

## 8.3 Interpretation of the corrections

The original supervised cache standardised complete MITDB recordings before window extraction. Its scale statistics therefore included the test region, and sensor chunks used a different scale. The corrected system filters and standardises each four-second model window and requires both neighbouring R peaks to be observed. Tests verify agreement between offline and streaming feature construction when the same peak locations are supplied.

The original 96.86% result is retained as a historical artifact, not the headline result of the corrected system. The preprocessing changes, removal of two edge beats per record, and refitting mean that the difference between the two runs is not a controlled estimate of one component's causal benefit.

The autoencoder remains the original frozen checkpoint. Its decoder resizes up-block tensors to their skip lengths and finally interpolates from 256 to 512 samples. The corrected architecture description reports those actual dimensions. Altering this trained architecture would require a separate training study.

## 8.4 Scientific limitations

1. Earlier labelled ECG from the same patients is available during supervised development. The result cannot establish cold-start accuracy.
2. The final segment was observed during earlier development. Correcting the preprocessing does not restore test independence. Selection optimism may remain, and neither the bootstrap nor the reported  $p$ -values removes it.
3. The bootstrap estimates subject variation with the fitted model fixed. It does not include retraining variability or extrapolation to a new cohort.
4. NSRDB is a nominally normal source cohort, with no beat-annotation screening in this preparation. Its ECG1 anatomical lead is unspecified. The two-epoch training record does not establish optimal convergence.
5. No controlled ablation quantifies the independent benefit of the four autoencoder residuals.
6. Output labels are project groups. AFib includes flutter on normal-like beat symbols; PVC includes ventricular escape beats; N includes atrial and junctional escape beats. No AF episode detection is evaluated.
7. Unsupported abnormal beats enter binary testing but are excluded from six-class scoring. The deployed gate has no validated unknown-class rejection.
8. Temporal and INCART scoring use annotated R peaks. Live peak detection introduces errors beyond those metrics. The prepared window needs roughly 2.44 seconds of post-R signal and the next observed beat, plus acquisition and computation delay.
9. INCART evaluation covers 5,416 beats but only one source patient and N/PVC reference support. More beats from this same patient cannot supply independent-patient evidence or validate absent classes.
10. The software requires the correct lead declaration, but cannot verify physical electrode placement. Signal checks reject non-finite/flat inputs and reported lead-off; they do not certify clinical signal quality.

## 8.5 Further research

The next experiment should preregister a patient-disjoint development, validation and test allocation; freeze preprocessing and thresholds before external scoring; and recruit enough independent subjects for every target class. PAC and external false alarms need particular attention. A matched ablation should compare the hierarchy with and without residual features, and prospective sensor testing should measure R-peak accuracy, classification errors, signal quality and end-to-end latency. The already viewed temporal and INCART records should be treated as development evidence in future work.

## 8.6 Final statement

This thesis delivers a corrected and reproducible single-lead research prototype with transparent retrospective results. It satisfies the measured aggregate targets within the stated evaluation, while exposing persistent PAC and external-domain weaknesses. It remains a research system, not a validated clinical diagnostic device [24].

# Bibliography

---

- [1] P. Kligfield, L. S. Gettes, J. J. Bailey, R. Childers, B. J. Deal, E. W. Hancock, G. van Herpen, J. A. Kors, P. Macfarlane, D. M. Mirvis, O. Pahlm, P. Rautaharju, and G. S. Wagner, “Recommendations for the standardization and interpretation of the electrocardiogram: Part i,” *Journal of the American College of Cardiology*, vol. 49, no. 10, pp. 1109–1127, 2007.  
→ [p1], [p23]
- [2] G. B. Moody and R. G. Mark, “The impact of the mit–bih arrhythmia database,” *IEEE Engineering in Medicine and Biology Magazine*, vol. 20, no. 3, pp. 45–50, 2001.  
→ [p1], [p6], [p11]
- [3] P. de Chazal, M. O’Dwyer, and R. B. Reilly, “Automatic classification of heartbeats using ecg morphology and heartbeat interval features,” *IEEE Transactions on Biomedical Engineering*, vol. 51, no. 7, pp. 1196–1206, 2004.  
→ [p1], [p6], [p9], [p21]
- [4] A. L. Goldberger, L. A. N. Amaral, L. Glass, J. M. Hausdorff, P. C. Ivanov, R. G. Mark, J. E. Mietus, G. B. Moody, C.-K. Peng, and H. E. Stanley, “Physiobank, physiotoolkit, and physionet: Components of a new research resource for complex physiologic signals,” *Circulation*, vol. 101, no. 23, pp. e215–e220, 2000.  
→ [p6], [p11]
- [5] G. B. Moody and R. G. Mark, “Mit–bih arrhythmia database.” <https://physionet.org/content/mitdb/1.0.0/>, 2005.  
→ [p6], [p11], [p13]
- [6] G. B. Moody, “Mit–bih normal sinus rhythm database.” <https://physionet.org/content/nsrdb/1.0.0/>, 1999.  
→ [p6], [p9], [p11], [p33]
- [7] S. Mousavi and F. Afghah, “Inter- and intra-patient ecg heartbeat classification for arrhythmia detection: A sequence to sequence deep learning approach,” in *IEEE International Conference on Acoustics, Speech and Signal Processing*, pp. 1308–1312, 2019.  
→ [p6], [p7], [p9], [p21]
- [8] J. Zhang, A. Liu, D. Liang, X. Chen, and M. Gao, “Interpatient ecg heartbeat classification with an adversarial convolutional neural network,” *Journal of Healthcare Engineering*, vol. 2021, p. 9946596, 2021.  
→ [p6], [p7], [p9], [p26]
- [9] H. Huang, J. Liu, Q. Zhu, R. Wang, and G. Hu, “Detection of inter-patient left and right bundle branch block heartbeats in ecg using ensemble classifiers,” *BioMedical Engineering OnLine*, vol. 13, p. 72, 2014.  
→ [p7], [p9]

- [10] S. Kiranyaz, T. Ince, and M. Gabbouj, “Real-time patient-specific ecg classification by 1-d convolutional neural networks,” *IEEE Transactions on Biomedical Engineering*, vol. 63, no. 3, pp. 664–675, 2016. → [p7]
- [11] U. R. Acharya, H. Fujita, O. S. Lih, Y. Hagiwara, J. H. Tan, and M. Adam, “Automated detection of arrhythmias using different intervals of tachycardia ecg segments with convolutional neural network,” *Information Sciences*, vol. 405, pp. 81–90, 2017. → [p7]
- [12] A. Y. Hannun, P. Rajpurkar, M. Haghpanahi, G. H. Tison, C. Bourn, M. P. Turakhia, and A. Y. Ng, “Cardiologist-level arrhythmia detection and classification in ambulatory electrocardiograms using a deep neural network,” *Nature Medicine*, vol. 25, pp. 65–69, 2019. → [p7]
- [13] A. H. Ribeiro, M. H. Ribeiro, G. M. M. Paixao, D. M. Oliveira, P. R. Gomes, J. A. Canazart, M. P. S. Ferreira, C. R. Andersson, P. W. Macfarlane, W. Meira, T. B. Schön, and A. L. P. Ribeiro, “Automatic diagnosis of the 12-lead ecg using a deep neural network,” *Nature Communications*, vol. 11, p. 1760, 2020. → [p7]
- [14] G. E. Hinton and R. R. Salakhutdinov, “Reducing the dimensionality of data with neural networks,” *Science*, vol. 313, no. 5786, pp. 504–507, 2006. → [p7], [p9], [p23]
- [15] P. Vincent, H. Larochelle, Y. Bengio, and P.-A. Manzagol, “Extracting and composing robust features with denoising autoencoders,” in *Proceedings of the 25th International Conference on Machine Learning*, pp. 1096–1103, 2008. → [p7], [p9], [p23], [p32]
- [16] V. Chandola, A. Banerjee, and V. Kumar, “Anomaly detection: A survey,” *ACM Computing Surveys*, vol. 41, no. 3, pp. 1–58, 2009. → [p7], [p24]
- [17] P. Malhotra, A. Ramakrishnan, G. Anand, L. Vig, P. Agarwal, and G. Shroff, “Lstm-based encoder–decoder for multi-sensor anomaly detection,” in *ICML 2016 Workshop on Anomaly Detection*, 2016. → [p7]
- [18] M. Thill, W. Konen, H. Wang, and T. Bäck, “Temporal convolutional autoencoder for unsupervised anomaly detection in time series,” *Applied Soft Computing*, vol. 112, p. 107751, 2021. → [p7], [p9]
- [19] J.-H. Jang, T. Y. Kim, H.-S. Lim, and D. Yoon, “Unsupervised feature learning for electrocardiogram data using the convolutional variational autoencoder,” *PLOS ONE*, vol. 16, no. 12, p. e0260612, 2021. → [p7]
- [20] L. Shan *et al.*, “Abnormal ecg detection based on an adversarial autoencoder,” *Frontiers in Physiology*, vol. 13, p. 961724, 2022. → [p7]
- [21] T. Saito and M. Rehmsmeier, “The precision–recall plot is more informative than the roc plot when evaluating binary classifiers on imbalanced datasets,” *PLOS ONE*, vol. 10, no. 3, p. e0118432, 2015. → [p8], [p9], [p27]

- [22] Y. Cui, M. Jia, T.-Y. Lin, Y. Song, and S. Belongie, "Class-balanced loss based on effective number of samples," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 9268–9277, 2019. → [p8], [p26]
- [23] C. M. Rutter, "Bootstrap estimation of diagnostic accuracy with patient-clustered data," *Academic Radiology*, vol. 7, no. 6, pp. 413–419, 2000. → [p8], [p10], [p27], [p45]
- [24] G. S. Collins, K. G. M. Moons, P. Dhiman, R. D. Riley, A. L. Beam, B. Van Calster, M. Ghassemi, X. Liu, J. B. Reitsma, M. van Smeden, A.-L. Boulesteix, J. C. Camaradou, L. A. Celi, S. Denaxas, A. K. Denniston, B. Glocker, R. M. Golub, H. Harvey, G. Heinze, M. M. Hoffman, A. P. Kengne, E. Lam, N. Lee, E. W. Loder, L. Maier-Hein, B. A. Mateen, M. D. McCradden, L. Oakden-Rayner, J. Ordish, R. Parnell, S. Rose, K. Singh, L. Wynants, and P. Logullo, "TRIPOD+AI statement: Updated guidance for reporting clinical prediction models that use regression or machine learning methods," *BMJ*, vol. 385, p. e078378, 2024. → [p8], [p10], [p57]
- [25] PhysioNet, "Frequently asked questions: What do the signal names mlii, v2, ... mean?," <https://archive.physionet.org/faq.shtml>. Accessed September 23, 2026. → [p13]
- [26] S. Butterworth, "On the theory of filter amplifiers," *Experimental Wireless and the Wireless Engineer*, vol. 7, pp. 536–541, 1930. → [p15], [p23]
- [27] P. Virtanen *et al.*, "Scipy 1.0: Fundamental algorithms for scientific computing in python," *Nature Methods*, vol. 17, pp. 261–272, 2020. → [p15]
- [28] P. Geurts, D. Ernst, and L. Wehenkel, "Extremely randomized trees," *Machine Learning*, vol. 63, pp. 3–42, 2006. → [p25]
- [29] E. Yakushenko, "St petersburg incart 12-lead arrhythmia database." <https://physionet.org/content/incartdb/1.0.0/>, 2008. → [p52]

# A Reproducibility Record

---

## A.1 Final experiment artifacts

The fixed temporal-holdout experiment is stored in:

- `experiments/corrected_temporal_holdout.json`;
- `experiments/corrected_temporal_holdout.joblib`; and
- `tools/final_temporal_holdout.py`.

The normal model is documented by `experiments/nsrdb_autoencoder_training.json`, produced by `tools/train_nsrdbe_autoencoder.py`. Its deployable checkpoint is `backend/app/storage/models/model_nsrdbe_primary_healthy.pt`. The deployable copy is registered by the backend model registry and stored in the backend model directory.

## A.2 Frozen constants

**Table A.1:** Settings used for the final experiment.

| Setting                             | Value                              |
|-------------------------------------|------------------------------------|
| Experiment seed                     | 20260801                           |
| Healthy-autoencoder seed            | 20260802                           |
| Healthy fitting/validation subjects | 15 / 3 NSRDB records               |
| Healthy fitting/validation windows  | 288,673 / 57,177                   |
| Bootstrap seed and replicates       | 20260801; 10,000                   |
| Development/test boundary           | 80% / 20% per record               |
| Candidate training/validation       | approximately 64% / 16% per record |
| Boundary gap                        | 16 beats before test               |
| Binary threshold                    | 0.5708333333                       |
| Binary ensemble                     | 240 fully grown Extra Trees        |
| Subtype ensemble                    | 260 fully grown Extra Trees        |
| Binary/subtype input sizes          | 242 / 242                          |
| Independent uncertainty units       | 45 subject clusters                |

## A.3 Saved success checks

The JSON artifact records three Boolean checks. Pooled whole-system accuracy above 90% is true. A whole-system subject-cluster lower 95% bound above 90% is true. Interval width at most three percentage points is true for the corrected rerun.

## A.4 Safety statement

The installed artifact is a research model. The final experiment uses patient overlap, one MITDB lead, and annotated R-peak locations during evaluation. It is not a certified medical device and must not be used to diagnose or treat a patient.

## A.5 Verification and historical record

The original August results remain in `experiments/final_temporal_holdout.json`. The corrected run saves all test predictions in `experiments/corrected_test_predictions.npz`, and independent checks are recorded in `experiments/final_verification.json`. The INCART result and predictions are `experiments/corrected_external_incart.json` and `experiments/corrected_external_predictions.npz`. Model SHA-256 digests link the evaluated bundle to its deployed copy.

From the repository root, use the backend Python environment to run `Final_Report_Research/tools/verify_final_results.py`. The corrected training entry point is `Final_Report_Research/tools/final_temporal_holdout.py`; it uses `tools/corrected_dataset.py` and the shared `backend/app/ml/beat_preparation.py`. The NSRDB checkpoint is frozen and is not retrained by this entry point. Running training again replaces the corrected artifacts, so preserve the current experiment directory before any future research change.